

EPIB 607: Inferential Statistics

Sahir Bhatnagar and James Hanley

2018-09-03

Contents

I	Preface	5
	Welcome	7
	Objectives	7
	Audience	7
	About these notes	7
	R Code Conventions	7
	Rendering Mathematical Formulae	8
	Development	8
	About the authors	8
	License	9
	Course Information	11
	Teaching strategy	11
	A focus on computation	11
	Grade Distribution	12
	Target Syllabus	13
	Descriptive Statistics	13
	Sampling Distributions	13
	Confidence Intervals	14
	Prerequisites	15
	Install R and RStudio	15
	R Packages	15
	Introduction to R	15
	Background Reading	15
II	Part I	17
1	Overview and Descriptive Statistics	19
	1.1 Populations and Samples	19
	1.2 Pictorial and Tabular Methods in Descriptive Statistics	21
	1.3 Measures of Location	21
	1.4 Measures of Variability	21
2	Probability	23
	Introduction	23
	2.1 Sample Spaces and Events	24
	2.2 Axioms, Interpretations, and Properties of Probability	26
	2.3 Counting Techniques	26
	2.4 Conditional Probability	27

2.5	Independence	31
3	Discrete Random Variables and Probability Distributions	33
	Introduction	33
3.1	Random Variables	34
3.2	Probability Distributions for Discrete Random Variables	35
3.3	Expected Values of Discrete Random Variables	36
3.4	Moments and Moment Generating Functions	36
3.5	The Binomial Probability Distribution	36
3.6	The Poisson Probability Distribution	36
4	Continuous Variables and Probability Distributions	37
	Introduction	37
A	Vectorization, *apply and for loops	39
A.1	Vectorization	39
A.2	Family of *apply functions	40
A.3	Creating dynamic documents with mapply	43
B	Appendix B	45

Part I

Preface

Welcome

Welcome to the course notes for [EPIB 607: Inferential Statistics](#) at McGill University.

Objectives

The aim of this course is to provide students with basic principles of statistical inference so that they can:

1. Understand the statistical methods section in a scientific paper.
2. Apply statistical methods in their own research.
3. Use the methods learned in this course as a foundation for more advanced biostatistics courses.

Audience

The principle audience is researchers in the natural and social sciences who have a basic understanding of differentiable and integral calculus, but haven't had an introductory course in statistics. This audience accepts that statistics has penetrated the life sciences pervasively and is required knowledge for both doing research and understanding scientific papers.

About these notes

These notes are a collection of useful links, videos, online resources and papers for an introductory course in statistics. The instructors have found that no single book sufficiently teaches all the topics covered in this course. Part of this is due to advancements in computing which have far outpaced the publication of modern textbooks. Indeed, the computer has replaced many of the calculations that were traditionally taught to be done by hand. We direct the readers to what we think is a good learning resource for a given topic (following the **Flipped Classroom** strategy). We also provide our own commentary and notes when we think its useful.

R Code Conventions

We use `R` code throughout these notes. When R code is displayed¹ it will be typeset using a monospace font with syntax highlighting enabled to ensure the differentiation of functions, variables, and so on. For example, the following adds 1 to 1

¹Reproduced with permission from <http://www.math.mcgill.ca/dstephens/>




Figure 1

```
a = 1L + 1L
a
```

Each code segment may contain actual output from R. Such output will appear in grey font prefixed by #>. For example, the output of the above code segment would look like so:

```
[1] 2
```

Rendering Mathematical Formulae

Throughout these notes, there will be mathematical symbols used to express the material. Depending on the version of the book, there are two different rendering engines.

- For the online version, the text uses [MathJax](#) to render mathematical notation for the web. In the event the formulae does not load for a specific chapter, first try to refresh the page. 9 times out of 10 the issue is related to the software library not loading quickly. You can also right-click to see the corresponding LaTeX code used to produce the equation.
- For the pdf version, the text is built using the recommended AMS LaTeX symbolic packages. As a result, there should be no issue displaying equations. An example of a mathematical rendering capabilities would be given as:

$$a^2 + b^2 = c^2$$

Development

This book is built with [bookdown](#) and is open source and freely available. This approach encourages contributions, ensures reproducibility and provides access to the material worldwide. The online version of the book is hosted at sahirbhatnagar.com/EPIB607 and kept up-to-date thanks to [Travis](#). The entire source code is available at <https://github.com/sahirbhatnagar/EPIB607>.

If you notice any errors, we would be grateful if you would let us know by filing an issue [here](#) or making a pull request by clicking the edit button in the top-left corner of the text:

The version of the book you are reading now was built on 2018-09-03 and was built on [Travis](#).

About the authors

Sahir Bhatnagar James Hanley

- Sahir R. Bhatnagar: Assistant Professor of Biostatistics - McGill University, Montreal, Canada.
 - Website: <https://sahirbhatnagar.com/>
 - Twitter: [syfi_24](#)
 - GitHub: <https://github.com/sahirbhatnagar>
- James A. Hanley: Professor of Biostatistics - McGill University, Montreal, Canada.
 - Webpage: <http://www.medicine.mcgill.ca/epidemiology/hanley/>

License

This work is licensed under a Creative Commons Attribution 4.0 International License

Course Information

- Instructor: [Sahir Bhatnagar](#)
- Teaching Assistants: [Kody Crowell](#), [Guanbo Wang](#), [Himasara Marasinghe](#)
- Website: <http://sahirbhatnagar.com/EPIB607/>
- Lectures: Monday 11:30am - 1:30pm, Thursday 8:30am - 10:30am
- Location: McMed 1034
- Office Hours: TBD
- Prerequisite(s): Calculus and Algebra
- Texts: *The Practice of Statistics in the Life Sciences*, 3rd Edition by Baldi & Moore.

Teaching strategy

This course will follow the **Flipped Classroom** model. Here, students are expected to have engaged with the material before coming to class (based on very precise pre-class instructions). The students will then be expected to answer a series of conceptual multiple choice questions using the [DALITE](#) online platform ([Bhatnagar et al., 2016](#)).

This allows the instructor to delegate the delivery of basic content and definitions to textbooks and videos, and enforces the idea that students cannot be simply passive recipients of information. This approach then allows the professor to focus valuable class time on nurturing efficient discussions surrounding the ideas within the content, guiding interactive exploration of typical misconceptions, and promoting collaborative problem solving with peers.

A focus on computation

Classic introductory statistics textbooks were written during a time when computers were still in their infancy. As such, even the newer editions heavily rely on *by-hand* computations such as looking up tables for tail probabilities. We take a modern approach and introduce computational methods in statistics with the statistical software program R.

Grade Distribution

Assignments	30%
DALITE Quizzes	10%
Midterm	20%
Project	10%
Final Exam	30%

Target Syllabus

```
knitr::kable(as.data.frame(rbind(c("JH", "James Hanley notes"),
                                c("AAO", "Against all odds video series"),
                                c("B&M", "The practice of statistics in the life sciences by Baldi and
                                c("Freedman", "Statistics by Freedman, Pisani, Purves, Adhikari"))),
  col.names = c("Abbreviation", "Description"),
  booktabs = TRUE)
```

Abbreviation	Description
JH	James Hanley notes
AAO	Against all odds video series
B&M	The practice of statistics in the life sciences by Baldi and Moore, 3rd edition
Freedman	Statistics by Freedman, Pisani, Purves, Adhikari

Descriptive Statistics

Topic	Video	Readings
Histograms	[link1](https://www.learner.org/courses/againstallodds/unitpages/unit03.html)	[pages 1-6](https://www.learner.org/resources/series/dataviz_chapter_7.html)
Density Plots	-	[dataviz chapter 7](https://serialsolutions.com/learning-center/online-learning/learners-series-dataviz-chapter-7)
Measures of Center	[link1](https://www.learner.org/courses/againstallodds/unitpages/unit04.html)	[pages 1-6](https://www.learner.org/resources/series/dataviz_chapter_7.html)
Boxplots	[link1](https://www.learner.org/courses/againstallodds/unitpages/unit05.html)	[pages 1-5](https://www.learner.org/resources/series/dataviz_chapter_7.html)
Standard Deviation	[link1](https://www.learner.org/courses/againstallodds/unitpages/unit06.html)	[pages 1-3](https://www.learner.org/resources/series/dataviz_chapter_7.html)
Data Visualization	[Hans Rosling BBC](https://www.youtube.com/watch?v=jbkSRLYSojo)	-

Sampling Distributions

Topic	Video
Parameters and Statistics	-
Sampling Distributions	[link1](https://www.learner.org/courses/againstallodds/unitpages/unit22.html)
Central Limit Theorem	[link1](https://www.learner.org/courses/againstallodds/unitpages/unit22.html), [link2](https://www.khanacademy.org/math/multivariable-calculus/multivariable-differentiation/differentiating-vector-valued-functions/a/multivariable-chain-rule-derivation/a190768c2)

Confidence Intervals

Topic	Video	Readings
Parameters and Statistics	–	[[JH notes]](https://www.dropbox.com/s/epgqkz3g0qklcp9/Ch14ConfidenceIntervalsJH2018.pdf?d)

Prerequisites

Install R and RStudio

All examples in this book are run in an [R](#) environment. You also need a recent version of [RStudio](#), which is a software application that facilitates how you interact with R. It is developed by data enthusiasts who consider statistics to be more than just simulations, formulas and proofs. RStudio emphasizes the following:

1. **Version control:** [Why I should use version control](#) especially for the [solo data analyst](#).
2. **Reproducible research:** seamless integration with [RMarkdown](#) for creating [dynamic documents](#) and presentations
3. **Creating R Packages:** seamless integration with the [devtools](#) package for creating software that implements your statistical method or analysis.

R Packages

The following packages will be called upon at some point, so please install them before getting started with the tutorials. Enter the following command in R:

```
install.packages(c("pacman", "knitr", "data.table", "rmarkdown", "tidyverse", "boot", "Hmisc"))
```

Introduction to R

Try out the interactive tutorial: <http://swirlstats.com/>

Background Reading

The greatest thing about R is that there are so many people out there willing to help you. R users are constantly writing tutorials and creating packages to make your analysis tasks easier. Here is a very targeted list that I suggest reading prior to starting the tutorials

1. [Writing Functions](#)
2. [for loops](#)
3. [apply vs. for](#)

Part II

Part I

Chapter 1

Overview and Descriptive Statistics¹

Statistical concepts and methods are not only useful but indeed often indispensable in understanding the world around us. They provide ways of gaining new insights into the behavior of many phenomena that you will encounter in your chosen field of specialization.

The discipline of statistics teaches us how to make intelligent judgments and informed decisions in the presence of *uncertainty and variation*.



Without uncertainty or variation, there would be little need for statistical methods or statisticians.

If the yield of a crop were the same in every field, if all individuals reacted the same way to a drug, if everyone gave the same response to an opinion survey, and so on, then a **single observation would reveal all desired information**.

1.1 Populations and Samples

We are constantly exposed to collections of facts, or data, both in our professional capacities and in everyday activities. The discipline of statistics provides methods for organizing and summarizing data and for drawing conclusions based on information contained in the data.

An investigation will typically focus on a well-defined collection of objects constituting a population of interest: - In one study, the population might consist of all gelatin capsules of a particular type produced during a specified period. - Another investigation might involve the population consisting of all individuals who received a B.S. in mathematics during the most recent academic year.

When desired information is available for all objects in the population, we have what is called a census. Constraints on time, money, and other scarce resources usually make a census impractical or infeasible. Instead, a subset of the population, **a sample**, is selected in some prescribed manner. Thus we might obtain a sample of pills from a particular production run as a basis for investigating whether pills are conforming to manufacturing specifications, or we might select a sample of last year's graduates to obtain feedback about the quality of the curriculum.

1.1.1 Variable

We are usually interested only in certain characteristics of the objects in a population: the amount of vitamin C in the pill, the gender of a mathematics graduate, the age at which the individual graduated, and so on. A variable is any

¹Devore and Berk.




Figure 1.1

characteristic whose value may change from one object to another in the population. Can be categorical (male/female) or numerical (temperature).

data type	description
univariate	consists of observations on a single variable
bivariate	observations are made on each of two variables
multivariate	more than two variables

1.1.2 Branches of Statistics

- **Descriptive Statistics:** summarize and describe important features of the data. Can be graphs (histograms, box-plots, and scatter plots), or numeric summaries (mean, standard deviations, and correlation coefficients)
- **Inferential Statistics:** Techniques for generalizing from a sample to a population. Having obtained a sample from a population, an investigator would frequently like to use sample information to draw some type of conclusion (make an inference of some sort) about the population. That is, the sample is a means to an end rather than an end in itself.



The focus of this course is Inferential statistics. But to get there we need to understand the basic concepts of probability

The relationship between the two disciplines can be summarized by saying that probability reasons from the population to the sample (**deductive reasoning**), whereas inferential statistics reasons from the sample to the population (**inductive reasoning**).

Before we can understand what a particular sample can tell us about the population, we should first understand the uncertainty associated with taking a sample from a given population. This is why we study probability before statistics.

Example 1.1 (Use of manual lap belts in cars equipped with automatic shoulder belt systems). *Probability:* assume that 50% of all drivers in a certain metropolitan area regularly use their lap belt → an assumption about the population. We might ask - How likely is it that a sample of 100 such drivers will include at least 70 who regularly use their lap belt? - How many of the drivers in a sample of size 100 can we expect to regularly use their lap belt?

Inference: a sample of 100 drivers of such cars revealed that 65 regularly use their lap belt. We might ask - Does this provide substantial evidence for concluding that more than 50% of all such drivers in this area regularly use their lap belt? We are attempting to use sample information to answer a question about the structure of the entire population from which the sample was selected.

1.2 Pictorial and Tabular Methods in Descriptive Statistics

1.3 Measures of Location

1.4 Measures of Variability


Chapter 2

Probability

Introduction¹

The random variation associated with *measurement* procedures in a scientific analysis requires a framework in which the **uncertainty** and **variability** that are inherent in the procedure can be handled. The key goal of Probability and Statistical modelling is to establish a mathematical framework within which *random* variation (due, for example, to experimental error or natural variation) can be quantified so that *systematic* variation (arising due to potentially important biological differences) can be studied.

Broadly, the *Scientific Process* involves several different stages:



Mathematical/Probabilistic modelling facilitates PREDICTION; *Statistical Analysis* provides the means of validation of predicted behaviour.

To explain the variation in observed data, we need to introduce the concept of a *probability distribution*. Essentially we need to be able to model, or specify, or compute the *chance* of observing the data that we collect or expect to collect. This will then allow us to assess how likely the data were to occur by chance alone, that is, how *surprising* the observed data are in light of an assumed theoretical model.

For example, consider two nucleotide sequences of the same length that we wish to assess for similarity:

Example 2.1 (Two nucleotide sequences).

Sequence 1	ATAGTAGATACGCACCGAGGA
Sequence 2	ATCTTAGATAGGCACTGAGGA

How can we assess sequence similarity formally? The number of discordant positions is 4, but how informative is that summary measure? Perhaps we need to assess the chance, for example, that a point mutation

$$A \rightarrow C$$

¹Reproduced with permission from <http://www.math.mcgill.ca/dstephens/>

occurs (as in the discordant position 3) in unit evolutionary time. Perhaps the chance of observing a sub-sequence

$$ATCTTA$$

rather than

$$ATAGTA$$

(in positions 1-6) is important.

- Is the hidden (or *latent*) structure in the sequence, corresponding to whether the sequence originates from a coding region or otherwise, important ?
- Can we even infer the hidden structure in light of the data we have observed ?

These questions can only really be answered when we have an understanding of randomness and variation. The framework that we will use to pose and answer such questions formally is given to us by *probability theory*.

Probability: A Measure of Uncertainty²

Often in life we are confronted by our own ignorance. Whether we are pondering tonight's traffic jam, tomorrow's weather, next week's stock prices, an upcoming election, or where we left our hat, often we do not know an outcome with certainty. Instead, we are forced to guess, to estimate, to hedge our bets.

Probability is the science of uncertainty.

It provides precise mathematical rules for understanding and analyzing our own ignorance. It does not tell us tomorrow's weather or next week's stock prices; rather, it gives us a **framework for working with our limited knowledge** and for **making sensible decisions based on what we do and do not know**.

To say there is a 40% chance of rain tomorrow is not to know tomorrow's weather. Rather, it is to **know what we do not know** about tomorrow's weather. In this course, we will develop a more precise understanding of what it means to say there is a 40% chance of rain tomorrow. We will learn how to work with ideas of randomness, probability, expected value, prediction, estimation, etc., in ways that are sensible and mathematically clear.

2.1 Sample Spaces and Events

2.1.1 Sample Spaces

Definition 2.1 (Sample Space). The sample space Ω is the set of possible outcomes of an experiment. Points ω in Ω are called sample outcomes, realizations, or elements.

Example 2.2 (Coin tossing). $\Omega = \{H, T\}$

Example 2.3 (Dice). $\Omega = \{1, 2, 3, 4, 5, 6\}$

Example 2.4 (Proportions). $\Omega = \{x : 0 \leq x \leq 1\}$

Example 2.5 (Time measurement). $\Omega = \{x : x > 0\} = \mathbb{R}^+$

Example 2.6 (Temperature measurement). $\Omega = \{x : a \leq x \leq b\} \subseteq \mathbb{R}$

²<http://www.utstat.toronto.edu/mikevans/jeffrosenthal/book.pdf>




Figure 2.1

Example 2.7 (Biological Sequence Analysis). The experiment may involve the observation of a nucleotide or protein sequence, so that the sample space Ω may comprise all sequences (of bases/amino acids) up to a given length, and a sample outcome would be a particular observed sequence.

There are two basic types of experiment: - Counting - Measurement

We shall see that these two types lead to two distinct ways of specifying probability distributions.

The collection of sample outcomes is a **set** (a collection of items) written as

$$s \in \Omega$$

if s is a member of the set Ω .

2.1.2 Events

Definition 2.2 (Event). An event E is a subset of the sample space Ω ($E \subseteq \Omega$). Events are usually denoted by upper case letters near the beginning of the alphabet, like A, B, C . An event which consists of only one outcome is called a simple (or elementary event); otherwise it is a compound event.



The sets Ω and E can be either be written as a list of items, for example,

$$E = \{s_1, s_2, \dots, s_n, \dots\}$$

which may a finite or infinite list, or can only be represented by a continuum of outcomes, for example

$$E = \{x : 0.6 < x \leq 2.3\}$$

Events are manipulated using **set theory** notation; if A and B are two events, $A, B \subseteq \Omega$, then

- $A \cup B$ is the set of outcomes that belong to A **or** to B , or to both,
- $A \cap B$ is the set of outcomes that belong to both A **and** to B .
- A^c (complement of A) is the set of outcomes **not** in A
- $A \setminus B = A \cap B^c$

The empty event will be denoted by \emptyset . Two events A and B are mutually exclusive if $A \cap B = \emptyset$, i.e., the collection of sample outcomes have no element in common.

2.2 Axioms, Interpretations, and Properties of Probability

Definition 2.3 (Axioms (basic properties) of Probability). To ensure that the probability assignments will be consistent with our intuitive notions of probability, all assignments should satisfy the following axioms (basic properties) of probability

- **AXIOM 1:** For any event A ,

$$P(A) \geq 0$$

- **AXIOM 2:**

$$P(\Omega) = 1$$

- **AXIOM 3:** If A_1, A_2, \dots is an infinite collection of disjoint events, then

$$P(A_1 \cup A_2 \cup \dots) = \sum_{i=1}^{\infty} P(A_i)$$

Proposition 2.1. $P(\emptyset) = 0$ where \emptyset is the null event. This in turn implies that the property contained in Axiom 3 is valid for a finite collection of events.

Proposition 2.2. For any event A ,

$$P(A) = 1 - P(A^c)$$

Proposition 2.3. For any event A ,

$$P(A) \leq 1$$

Proposition 2.4. For any events A and B ,

$$P(A \cup B) = P(A) + P(B) - P(A \cap B)$$

2.3 Counting Techniques

When the various outcomes of an experiment are equally likely (the same probability is assigned to each simple event), the task of computing probabilities reduces to counting. In particular, if N is the number of outcomes in a sample space and $N(A)$ is the number of outcomes contained in an event A , then

$$P(A) = \frac{N(A)}{N}$$

Proposition 2.5 (Product rule for ordered pairs). If the first element or object of an ordered pair can be selected in n_1 ways, and for each of these n_1 ways the second element of the pair can be selected in n_2 ways, then the number of pairs is $n_1 \cdot n_2$.

2.3.1 Permutations

Definition 2.4 (Permutation). Any ordered sequence of k objects taken from a set of n distinct objects is called a permutation of size k of the objects. The number of permutations of size k that can be constructed from the n objects is denoted by P_k^n :

$$P_k^n = \frac{n!}{(n-k)!}$$

2.3.2 Combinations

Definition 2.5 (Combination). Given a set of n distinct objects, any unordered subset of size k of the objects is called a combination. The number of combinations n of size k that can be formed from n distinct objects will be denoted by

$$\binom{n}{k} = \frac{n!}{k!(n-k)!} = \frac{P_k^n}{k!}$$

2.4 Conditional Probability

Conditional probability is the means by which probabilities are updated in the light of new information. We examine how the information *an event B has occurred* affects the probability assigned to A .

Example 2.8 (Flipping Coins). We flip three different fair coins, and

$$\Omega = HHH, HHT, HTH, HTT, THH, THT, TTH, TTT$$

with $P(s) = 1/8$ for each $s \in \Omega$. What is the probability that the first coin comes up heads?

$$P(\text{first coin heads}) = P(HHH, HHT, HTH, HTT) = 4/8 = 1/2$$

But suppose now that an informant tells us that exactly two of the three coins came up heads. Now what is the probability that the first coin was heads? if we know that exactly two of the coins were heads, then we know that the outcome was one of HHT, HTH, THH .

Because those three outcomes should (in this case) still all be equally likely, and because only the first two correspond to the first coin being heads, we conclude the following: If we know that exactly two of the three coins are heads, then the probability that the first coin is heads is $2/3$.

More precisely, we have computed a conditional probability. That is, we have determined that, conditional on knowing that exactly two coins came up heads, the conditional probability of the first coin being heads is $2/3$. We write this in mathematical notation as

$$P(\text{first coin heads} | \text{two coins heads}) = 2/3.$$

Here the vertical bar $|$ stands for *conditional on* or *given that*.

Example 2.9 (Assembly Lines). Complex components are assembled in a plant that uses two different assembly lines, A and A^c . Line A uses older equipment than A^c , so it is somewhat slower and less reliable. B are the defective components and B^c are the nondefective.

Line	Condition	
	B	B^c
A	2	6
A^c	1	9

The sales manager randomly selects 1 of these 18 components for a demonstration

$$P(\text{line A component was selected}) = P(A) = \frac{N(A)}{N} = \frac{8}{18} = 0.444$$

However, if the chosen component turns out to be defective, then the event B has occurred, so the component must have

been 1 of the 3 in the B column of the table. Since these 3 components are equally likely among themselves after B has occurred

$$P(\text{line A component was selected}|\text{Defective}) = \frac{2}{3} = \frac{2/18}{3/18} = \frac{P(A \cap B)}{P(B)}$$

In Example 2.9, the conditional probability is expressed as a ratio of **unconditional probabilities**. The numerator is the probability of the intersection of the two events, whereas the denominator is the probability of the conditioning event B . Given that B has occurred, the relevant sample space is no longer Ω but consists of just outcomes in B ; A has occurred if and only if *one of the outcomes in the intersection* occurred, so the conditional probability of A given B is proportional to $P(A \cap B)$. The proportionality constant $1/P(B)$ is used to ensure that the probability $P(B|B)$ of the new sample space B equals 1.

Definition 2.6 (Conditional Probability). Given two events A and B , with $P(B) > 0$, the conditional probability of A given B is equal to

$$P(A|B) = \frac{P(A \cap B)}{P(B)}$$

In example 2.8, let

- $A = HHH, HHT, HTH, HTT$ be the event that the first coin is heads

- $B = HHT, HTH, THH$ be the event that exactly two coins were heads

It follows that

$$A \cap B = HHT, HTH$$

Therefore

$$P(A|B) = \frac{P(A \cap B)}{P(B)} = \frac{P(HHT, HTH)}{P(HHT, HTH, THH)} = \frac{2/8}{3/8} = \frac{2}{3}$$

Example 2.10 (Balanced die). Suppose a balanced die is tossed in the next room. We are told that a number less than 4 was observed. What is the probability the number was either 1 or 2?

Example 2.11 (Two Balanced Dice v1). Toss two balanced dice. Let $A = \{\text{sum of 5}\}$ and $B = \{\text{first die is } \leq 2\}$. Find $P(A|B)$

Example 2.12 (Two Balanced Dice v2). Two balanced dice are tossed. What is the probability that the first die gives a number less than three, given that the sum is odd?

Example 2.13 (Unbalanced Die). Toss an unbalanced die with probs $P(1) = .1, P(2) = .1, P(3) = .3, P(4) = .2, P(5) = .1, P(6) = .2$. Let $A = \geq 5$ and $B = \geq 2$. Find $P(A|B)$.

Example 2.14 (Two Balanced Coins). Two balanced coins were tossed, and it is known that at least one was a head. What is the probability that both were heads?

Example 2.15 (Two Cards). Two cards are drawn without replacement from a standard deck. Find the probability that

- 1) the second is an ace, given that the first is not an ace.
- 2) the second is an ace.
- 3) the first was an ace, given that the second is an ace.

Example 2.16 (Numbers in a Hat). The numbers 1 to 5 are written on five slips of paper and placed in a hat. Two slips are drawn at random without replacement. What is the probability that the first number is 3, given a sum of seven?

Example 2.17 (One Card). A card is selected at random (i.e. every card has the same probability of being chosen) from a deck of 52. What is the probability it is a red card or a face card?

Definition 2.6 immediately leads to the *multiplication formula*

Definition 2.7 (Multiplicative Rule).

$$P(A \cap B) = P(A|B)P(B)$$

and

$$P(A \cap B) = P(B|A)P(A)$$

This allows us to compute the joint probability of A and B when we are given the probability of B and the conditional probability of A given B , and vice versa.

Example 2.18 (Fish in a Tank). A tank has three red fish and two blue fish. Two fish are chosen at random and without replacement. What is the probability of getting

- 1) red fish first and then a blue fish?
- 2) both fish red?
- 3) one red fish and one blue fish?

2.4.1 Law of Total Probability

Recall that events A_1, A_2, \dots, A_k are mutually exclusive if no two have any common outcomes. The events are exhaustive if one A_i must occur, so that

$$A_1 \cup A_2 \cup \dots \cup A_k = \Omega$$

Theorem 2.1 (Law of Total Probability). Let A_1, A_2, \dots, A_k be mutually exclusive and exhaustive events. Then for any other event B ,

$$P(B) = P(B|A_1)P(A_1) + \dots + P(B|A_k)P(A_k) = \sum_{i=1}^k P(B|A_i)P(A_i)$$

Proof: Because the A_i 's are mutually exclusive and exhaustive, if B occurs it must be in conjunction with exactly one of the A_i 's. That is, $B = (A_1 \text{ and } B) \text{ or } \dots \text{ or } (A_k \text{ and } B)$ which is equal to $(A_1 \cap B) \cup \dots \cup (A_k \cap B)$, where the events $(A_i \cap B)$ are mutually exclusive.

Thus we have

$$P(B) = \sum_{i=1}^k P(A_i \cap B) = \sum_{i=1}^k P(B|A_i)P(A_i)$$

Example 2.19 (Long Hair). Suppose a class contains 60% girls and 40% boys. Suppose that 30% of the girls have long hair, and 20% of the boys have long hair. A student is chosen uniformly at random from the class. What is the probability that the chosen student will have long hair?




Figure 2.2

2.4.2 Bayes' Rule

Theorem 2.2 (Bayes' Rule). Let A_1, \dots, A_k be a collection of mutually exclusive and exhaustive events with $P(A_i) > 0$ for $i = 1, \dots, k$. Then for any other event B , for which $P(B) > 0$, we have

$$P(A_j|B) = \frac{P(A_j \cap B)}{P(B)} = \frac{P(B|A_j)P(A_j)}{\sum_{i=1}^k P(B|A_i)P(A_i)}, \quad j = 1, \dots, k$$

The transition from the second to the third expression in Theorem 2.2 rests on using the multiplication rule in the numerator and the law of total probability in the denominator.

Example 2.20 (Urns). Suppose urn #1 has 3 red and 2 blue balls, and urn #2 has 4 red and 7 blue balls. Suppose one of the two urns is selected with probability $1/2$ each, and then one of the balls within that urn is picked uniformly at random.

- 1) What is the probability that urn #2 is selected at the first stage (event A) and a blue ball is selected at the second stage (event B)?
- 2) Compute the probability that a blue ball is obtained.
- 3) Now suppose we are given the information that the ball picked is blue. What is probability that we had selected urn #2?

Example 2.21 (Large Bridges). There are three Canadian firms which build large bridges, firm 1, firm 2, and firm 3. 20% of Canadian large bridges have been built by firm 1, 30% by firm 2, and the rest by firm 3. 5% of the bridges built by firm 1 have collapsed, while 10% of those by firm 2 have collapsed, and 30% by firm 3 have collapsed.

- 1) What is the probability that a bridge collapses?

- 2) Suppose it is reported in tomorrow's newspaper that a large bridge has collapsed. What is the probability it was built by firm 1?

2.5 Independence

If we flip a fair coin twice, then the probability of two heads is $1/2 \times 1/2$. We multiply the probabilities because we regard the two tosses as independent. The formal definition of independence is as follows:

Definition 2.8 (Independent Events). Two events A and B are independent if

$$P(A \cap B) = P(A)P(B)$$

Now, because $P(A|B) = P(A \cap B)/P(B)$, we see that A and B are independent if and only if $P(A|B) = P(A)$ or $P(B|A) = P(B)$, provided that $P(A) > 0$ and $P(B) > 0$. Definition 2.8 has the advantage that it remains valid even if $P(B) = 0$ or $P(A) = 0$, respectively. Intuitively, events A and B are independent if neither one has any impact on the probability of the other.

Example 2.22 (Toss a fair coin 10 times). Toss a fair coin 10 times. Let A = at least one head. Let T_j be the event that tails occurs on the j^{th} toss. Find $P(A)$

Example 2.23 (Unbalance Die Revisited). In Example 2.13, if A is the event that the die was 5, and B is the event that the coin was tails, then calculate $P(A)$, $P(B)$ and $P(A \cap B)$

Chapter 3

Discrete Random Variables and Probability Distributions

Introduction

In Chapter 2, we discussed the probability model as the central object of study in the theory of probability. This required defining a probability measure P on a class of subsets of the sample space Ω . For example, for an experiment with possible sample outcomes denoted by the *sample space* Ω , an *event* E was defined as any collection of sample outcomes, that is, any subset of the set Ω .

$$\begin{array}{ccccccc} \text{EXPERIMENT} & \longrightarrow & \text{SAMPLE OUTCOMES} & \longrightarrow & \text{EVENTS} & \longrightarrow & \text{PROBABILITIES} \\ & & \Omega = \{s_1, s_2, \dots\} & & \longrightarrow E \subseteq S & & \longrightarrow P(E) \end{array}$$

In this framework, it is necessary to consider each experiment with its associated sample space separately - the nature of sample space Ω is typically different for different experiments.

Example 3.1 (Rainy days). Count the number of days in February which have zero precipitation.

$$\Omega = \{0, 1, 2, \dots, 28\}$$

Let $E_i = i$ days have zero precipitation. E_0, \dots, E_{28} partition Ω .

Example 3.2 (Football Match). Count the number of goals in a football match.

$$\Omega = \{0, 1, 2, 3, \dots\}$$

Let $E_i = i$ goals in the match. E_0, E_1, E_2, \dots partition Ω



In both of these examples, we need a formula to specify each

$$P(E_i) = p_i$$

Example 3.3 (Operating Temperature). Measure the operating temperature of an experimental process.

$$\Omega = \{x : x > T_{min}\}$$

Here it is difficult to express

$$P(\text{"Measurement is } x)$$

but possible to think about

$$P(\text{"Measurement is } \leq x) = F(x)$$

and now we seek a formula for $F(x)$ which is a simpler way of presenting a particular probability assignment.

This chapter is concerned with the definitions of random variables, distribution functions $F(x)$, probability/density functions $f(x)$, and the development of the concepts necessary for carrying out calculations for a probability model using these entities (?).

The concept of a random variable allows us to pass from the experimental outcomes themselves to a numerical function of the outcomes. There are two fundamentally different types of random variables (?):

- i) discrete random variables
- ii) continuous random variables

In this chapter, we examine the basic properties and discuss the most important examples of **discrete** variables. Chapter 4 focuses on continuous random variables.

3.1 Random Variables

A general notation useful for all such examples can be obtained by considering a sample space that is **equivalent** to Ω for a general experiment, but whose form is more familiar.

Definition 3.1 (Random Variable). A random variable X on Ω is a function from the sample space Ω to the set \mathbb{R} of all real numbers denoted by

$$X : \Omega \rightarrow \mathbb{R}$$

Let R_X denote the range of X .

X is called a *discrete* random variable if R_X is a countable set.



Random variables are customarily denoted by uppercase letters, such as X and Y , lowercase letters to represent some particular value of the corresponding random variable. The notation $X(s) = x$ means that x is the value associated with the outcome s by the rv X .

Example 3.4 (Coin Toss). Suppose a coin is tossed three times. Let X be the number of heads observed. The sample space is

$$\Omega = \left\{ \underbrace{HHT}_3, \underbrace{HHT}_2, \underbrace{HTH}_2, \underbrace{HTT}_1, \underbrace{THH}_2, \underbrace{THT}_1, \underbrace{TTH}_1, \underbrace{TTT}_0 \right\}$$




Figure 3.1 A random variable

Figure 3.1

That is, we have $X(HHH) = 3$, $X(HHT) = 2$, $X(HTH) = 2$, and so on. Hence $R_X = \{0, 1, 2, 3\}$

Example 3.5 (A Very Simple Random Variable). Let the random variable $X : \{\text{rain, snow, clear}\} \rightarrow \mathbb{R}$ by $X(\text{rain}) = 3$, $X(\text{snow}) = 6$, and $X(\text{clear}) = -2.7$.

We now present several further examples. The point is, we can define random variables any way we like, as long as they are functions from the sample space to \mathbb{R} .

Example 3.6 (A Very Simple Random Variable 2). For the case $\Omega = \{\text{rain, snow, clear}\}$, we might define a second random variable Y by saying that $Y = 0$ if it rains, $Y = -1/2$ if it snows, and $Y = 7/8$ if it is clear. That is $Y(\text{rain}) = 0$, $Y(\text{snow}) = 1/2$, and $Y(\text{rain}) = 7/8$.

Example 3.7 (A Very Simple Random Variable 3). If the sample space corresponds to flipping three different coins, then we could let X be the total number of heads showing, let Y be the total number of tails showing, let $Z = 0$ if there is exactly one head, and otherwise $Z = 17$.

Example 3.8 (Constants as Random Variables). As a special case, every constant value c is also a random variable, by saying that $c(s) = c$ for all $s \in \Omega$. Thus, 5 is a random variable, as is 3 or -21.6.

Example 3.9 (Indicator Functions). If A is any event, then we can define the indicator function of A , written I_A , to be the random variable

$$I_A(s) = \begin{cases} 1 & s \in A \\ 0 & s \notin A \end{cases}$$

Suppose X is a random variable. We know that different states s occur with different probabilities. It follows that $X(s)$ also takes different values with different probabilities. These probabilities are called the **distribution** of X ; we consider them next.

3.2 Probability Distributions for Discrete Random Variables

Because random variables are defined to be functions of the outcome s , and because the outcome s is assumed to be random (i.e., to take on different values with different probabilities), it follows that the value of a random variable will itself be random (as the name implies).

Specifically, if X is a random variable, then what is the probability that X will equal some particular value x ? Well, $X = x$ precisely when the outcome s is chosen such that $X(s) = x$.

3.3 Expected Values of Discrete Random Variables

3.4 Moments and Moment Generating Functions

3.5 The Binomial Probability Distribution

3.6 The Poisson Probability Distribution

Chapter 4

Continuous Variables and Probability Distributions

Introduction

Appendix A

Vectorization, *apply and for loops

This section will cover the basics of vectorizations, the `*apply` family of functions and `for` loops.

A.1 Vectorization

Almost everything in R is a vector. A scalar is really a vector of length 1 and a `data.frame` is a collection of vectors. An nice feature of R is its vectorized capabilities. Vectorization indicates that a function operates on a whole vector of values at the same time and not just on a single value¹. If you have ever taken a basic linear algebra course, this concept will be familiar to you.

Take for example two vectors:

$$\begin{bmatrix} 1 \\ 2 \\ 3 \end{bmatrix} + \begin{bmatrix} 1 \\ 2 \\ 3 \end{bmatrix} = \begin{bmatrix} 2 \\ 4 \\ 6 \end{bmatrix}$$

The corresponding R code is given by:

```
a <- c(1,2,3)
b <- c(1,2,3)
a+b
```

```
## [1] 2 4 6
```

Many of the base functions in R are already vectorized. Here are some common examples:

```
# generate a sequence of numbers from 1 to 10
(a <- 1:10)
```

```
## [1] 1 2 3 4 5 6 7 8 9 10
```

```
# sum the numbers from 1 to 10
sum(a)
```

¹<http://www.dummies.com/how-to/content/how-to-vectorize-your-functions-in-r.html>

```
## [1] 55
```

```
# calculate sums of each column
colSums(iris[, -5])
```

```
## Sepal.Length Sepal.Width Petal.Length Petal.Width
##      876.5      458.6      563.7      179.9
```

Exercise: What happens when you sum two vectors of different lengths?

A.2 Family of *apply functions

- apply, lapply and sapply are some of the most commonly used class of functions in R
- *apply functions are not necessarily faster than loops, but can be easier to read (and vice versa)
- apply is used when you need to perform an operation on every row or column of a matrix or data.frame
- lapply and sapply differ in the format of the output. The former returns a list while the latter returns a vector
- There are other *apply functions such as tapply, vapply and mapply with similar functionality and purpose

A.2.1 Loops vs. Apply

```
# Getting the row means of two columns
# Generate data
N <- 10000
x1 <- runif(N)
x2 <- runif(N)
d <- as.data.frame(cbind(x1, x2))
head(d)
```

```
##      x1      x2
## 1 0.57632 0.9615
## 2 0.56474 0.1950
## 3 0.07399 0.3001
## 4 0.45387 0.3823
## 5 0.37328 0.1197
## 6 0.33132 0.9891
```

```
# Loop:
# create a vector to store the results in
rowMeanFor <- vector("double", N)

for (i in seq_len(N)) {
  rowMeanFor[[i]] <- mean(c(d[i, 1], d[i, 2]))
}
```

```
# Apply:
rowMeanApply <- apply(d, 1, mean)
```

```
# are the results equal
all.equal(rowMeanFor, rowMeanApply)
```



```
## [1] TRUE
```

A.2.2 Descriptive Statistics using *apply

```
data(women)
# data structure
str(women)
```

```
## 'data.frame': 15 obs. of 2 variables:
## $ height: num 58 59 60 61 62 63 64 65 66 67 ...
## $ weight: num 115 117 120 123 126 129 132 135 139 142 ...
```

```
# calculate the mean for each column
apply(women, 2, mean)
```

```
## height weight
## 65.0 136.7
```

```
# apply 'fivenum' function to each column
vapply(women, fivenum, c("Min." = 0, "1st Qu." = 0, "Median" = 0,
                        "3rd Qu." = 0, "Max." = 0))
```

```
##           height weight
## Min.      58.0 115.0
## 1st Qu.   61.5 124.5
## Median    65.0 135.0
## 3rd Qu.   68.5 148.0
## Max.      72.0 164.0
```

A.2.3 Creating new columns using sapply

You can apply a *user defined function* to columns or the entire data frame:

```
# the output of sapply is a vector
# the 's' in sapply stands for 'simplified' apply
mtcars$gear2 <- sapply(mtcars$gear,
                      function(i) if (i==4) "alot" else "some")

head(mtcars)[,c("gear", "gear2")]
```

```
##           gear gear2
## Mazda RX4      4  alot
## Mazda RX4 Wag  4  alot
## Datsun 710      4  alot
## Hornet 4 Drive  3  some
## Hornet Sportabout 3  some
## Valiant        3  some
```

A.2.4 Applying functions to subsets using `tapply`

```
# Fisher's famous dataset
data(iris)
str(iris)

## 'data.frame':   150 obs. of  5 variables:
## $ Sepal.Length: num  5.1 4.9 4.7 4.6 5 5.4 4.6 5 4.4 4.9 ...
## $ Sepal.Width : num  3.5 3 3.2 3.1 3.6 3.9 3.4 3.4 2.9 3.1 ...
## $ Petal.Length: num  1.4 1.4 1.3 1.5 1.4 1.7 1.4 1.5 1.4 1.5 ...
## $ Petal.Width : num  0.2 0.2 0.2 0.2 0.2 0.4 0.3 0.2 0.2 0.1 ...
## $ Species      : Factor w/ 3 levels "setosa","versicolor",...: 1 1 1 1 1 1 1 1 1 1 ...

# mean sepal length by species
tapply(iris$Sepal.Length, iris$Species, mean)

##      setosa versicolor virginica
##      5.006      5.936      6.588
```

A.2.5 Nested for loops using `mapply`

`mapply` is my favorite base R function and here are some reasons why:

- Using `mapply` is equivalent to writing nested `for` loops except that it is 100% more human readable and less prone to errors
- It is an effective way of conducting simulations because it iterates of many arguments

Let's say you want to generate random samples from a normal distribution with varying means and standard deviations. Of course the brute force way would be to write out the command once, copy paste as many times as you want, and then manually change the arguments for `mean` and `sd` in the `rnorm` function as so:

```
v1 <- rnorm(100, mean = 5, sd = 1)
v2 <- rnorm(100, mean = 10, sd = 5)
v3 <- rnorm(100, mean = -3, sd = 10)
```

This isn't too bad for three vectors. But what if you want to generate many more combinations of means and sds? Furthermore, how can you keep track of the parameters you used? Now let's consider the `mapply` function:

```
means <- c(5,10,-3) ; sds <- c(1,5,10)

# MoreArgs is a list of arguments that dont change
randomNormals <- mapply(rnorm, mean = means, sd = sds,
                        MoreArgs = list(n = 100))

head(randomNormals)

##      [,1]  [,2]  [,3]
## [1,] 3.836  8.771  5.144
## [2,] 4.525  9.376 -2.280
## [3,] 4.072 13.144  2.940
## [4,] 4.737 18.210 -13.118
```



Figure A.1

```
## [5,] 5.690 22.951 -7.008
## [6,] 4.826 7.615 -15.323
```

The following diagram (from [r4ds](#)) describes exactly what is going on in the above function call to `map2`:

Advantages:

1. Result is automatically stored in a matrix
2. The parameters are also saved in R objects so that they can be easily manipulated and/or recovered

Consider a more complex scenario where you want to consider many possible combinations of means and sds. We take advantage of the `expand.grid` function to create a `data.frame` of simulation parameters:

```
simParams <- expand.grid(means = 1:10,
                        sds = 1:10)

randomNormals <- map2(rnorm, mean = simParams$means,
                     sd = simParams$sds,
                     MoreArgs = list(n = 100))

dim(randomNormals)

## [1] 100 100
```

A.3 Creating dynamic documents with `mapply`

`mapply` together with the `rmarkdown` package (Allaire et al., 2018) can be very useful to create dynamic documents for exploratory analysis. We illustrate this using the Motor Trend Car Road Tests data which comes pre-loaded in R.

The data was extracted from the 1974 Motor Trend US magazine, and comprises fuel consumption and 10 aspects of automobile design and performance for 32 automobiles (1973–74 models).

Copy the code below in a file called `mapplyRmarkdown.Rmd`:

Copy the code below in a file called `boxplotTemplate`:

Appendix B

Appendix B

Bibliography

Allaire, J., Xie, Y., McPherson, J., Luraschi, J., Ushey, K., Atkins, A., Wickham, H., Cheng, J., and Chang, W. (2018). *rmarkdown: Dynamic Documents for R*. R package version 1.10.

Bhatnagar, S., Lasry, N., Desmarais, M., and Charles, E. (2016). Dalite: Asynchronous peer instruction for moocs. In *European Conference on Technology Enhanced Learning*, pages 505–508. Springer.