Sampling Distributions, the Central Limit Theorem (CLT) and Confidence Intervals

Sahir Bhatnagar and James Hanley

EPIB 607
Department of Epidemiology, Biostatistics, and Occupational Health
McGill University

sahir.bhatnagar@mcgill.ca
https://sahirbhatnagar.com/EPIB607/

November 3, 2018

Parameters, Samples, and Statistics

Parameters, Samples, and Statistics

- **Paramter**: An unknown numerical constant pertaining to a population/universe, or in a statistical model.
  - μ: population mean π: population proportion
- **Statistic**: A numerical quantity calculated from a sample. The empirical counterpart of the parameter, used to *estimate* it.
  - $\bar{y}$: sample mean p: sample proportion

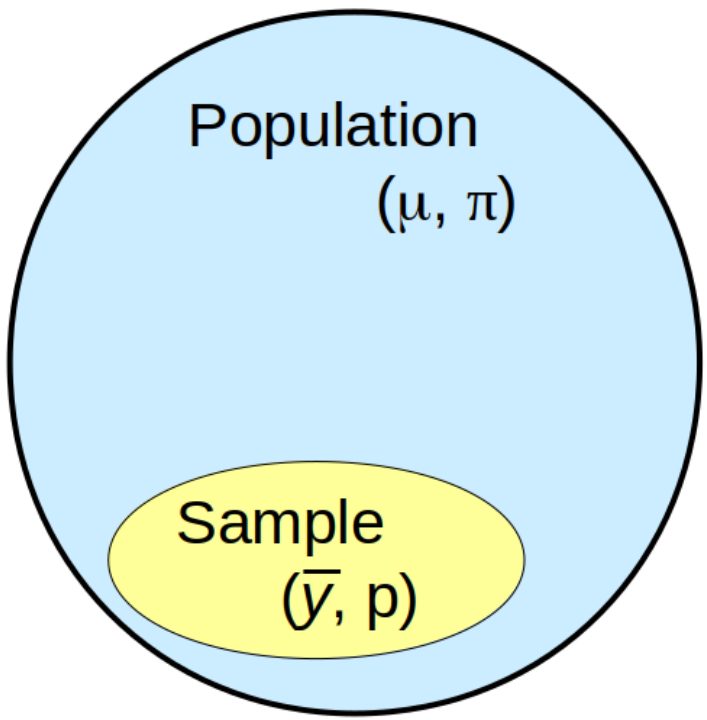

Examples

Proportions:

- Proportion of Earth's surface covered by water
- Proportion who saw a medical doctor last year
- Proportion of Québécois who don't have a family doctor

Means:

- Mean depth in n randomly selected ocean locations
- Mean household size in n randomly selected households.
- Median number of persons under-5 in a sample of n households

Samples must be random

- The validity of inference will depend on the way that the sample was collected. If a sample was collected badly, no amount of statistical sophistication can rescue the study.

- Samples should be **random**. That is, there should be no systematic set of characteristics that is related to the scientific question of interest that causes some people to be more likely to be sampled than others. The simplest type of randomization selects members from the population with equal probability (a uniform distribution).

- When conducting a study, it is always better to seek statistical advice sooner rather than later. Get a statistician involved at the *planning* stage of the study... by the analysis stage, it may be too late!

Samples must be random - No cheating!

Do not cheat by

- Taking 5 people from the *same* household to estimate
  - proportion of Québécois who don't have a family doctor
  - who saw a medical doctor last year
  - average rent

- Sampling the depth of the ocean *only around Montreal* to estimate
  - proportion of Earth's surface covered by water

Collecting data takes effort

In general

- The larger the sample $\rightarrow$ the more accurate the estimate (if sampling is done correctly)

CAVEAT

- Collecting more data takes effort and money!
- We will also soon discover the curse of the $\sqrt{n}$

Collecting data takes effort

In general

- The larger the sample $\rightarrow$ the more accurate the estimate (if sampling is done correctly)

CAVEAT

- Collecting more data takes effort and money!
- We will also soon discover the curse of the $\sqrt{n}$

Sampling Distributions

Sampling Distributions

- Given a sample of n observations from a population, we will be calculating estimates of the population mean, proportion, standard deviation, and various other population characteristics (parameters)
- Prior to obtaining data, there is uncertainty as to which of all possible samples will occur
- Because of this, estimates such as $\bar{y}$ (the sample mean) will vary from one sample to another

Sampling Distributions

- The behavior of such estimates in many samples of equal size is described by what are called **sampling distributions**

- B&M definition: The sampling distribution of a statistic is the distribution of values taken by the statistic in all possible samples of the same size from the same population.

Why are sampling distributions important?

- They tell us how far from the target (true value of the parameter) our statistical *shot* at it (i.e. the statistic calculated form a sample) is likely to be, or, to have been.

- Thus, they are used in confidence intervals for parameters. Specific sampling distributions (based on a null value for the parameter) are also used in statistical tests of hypotheses.

Exercise 1: How Deep is the Ocean?

- We will get a sense of what a sampling distribution is in Exercise 1
- **CAVEAT**: This is a luxury using a toy example. In actual studies, we only get one shot!

Exercise 1 Results

Sampling distribution: proportion covered by water

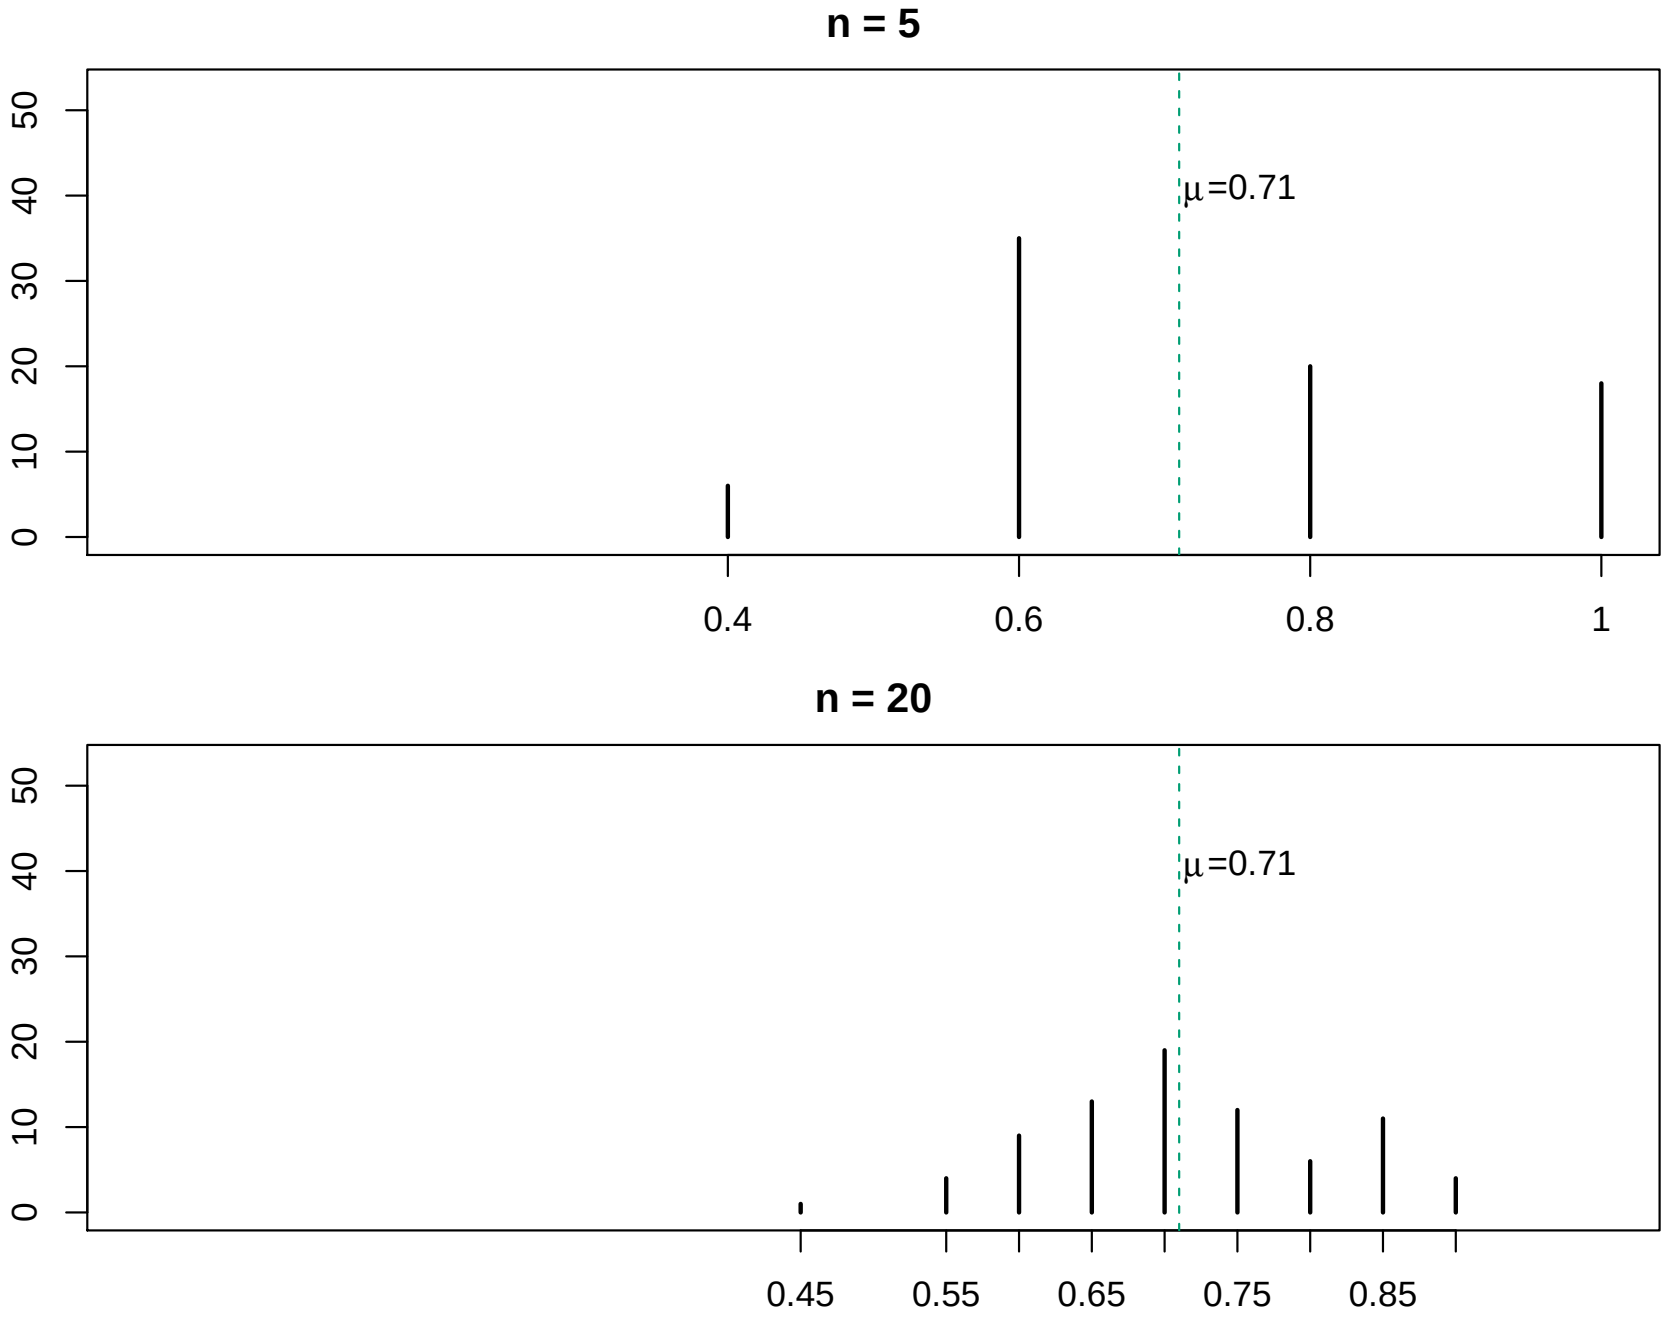

Sampling distribution: proportion covered by water

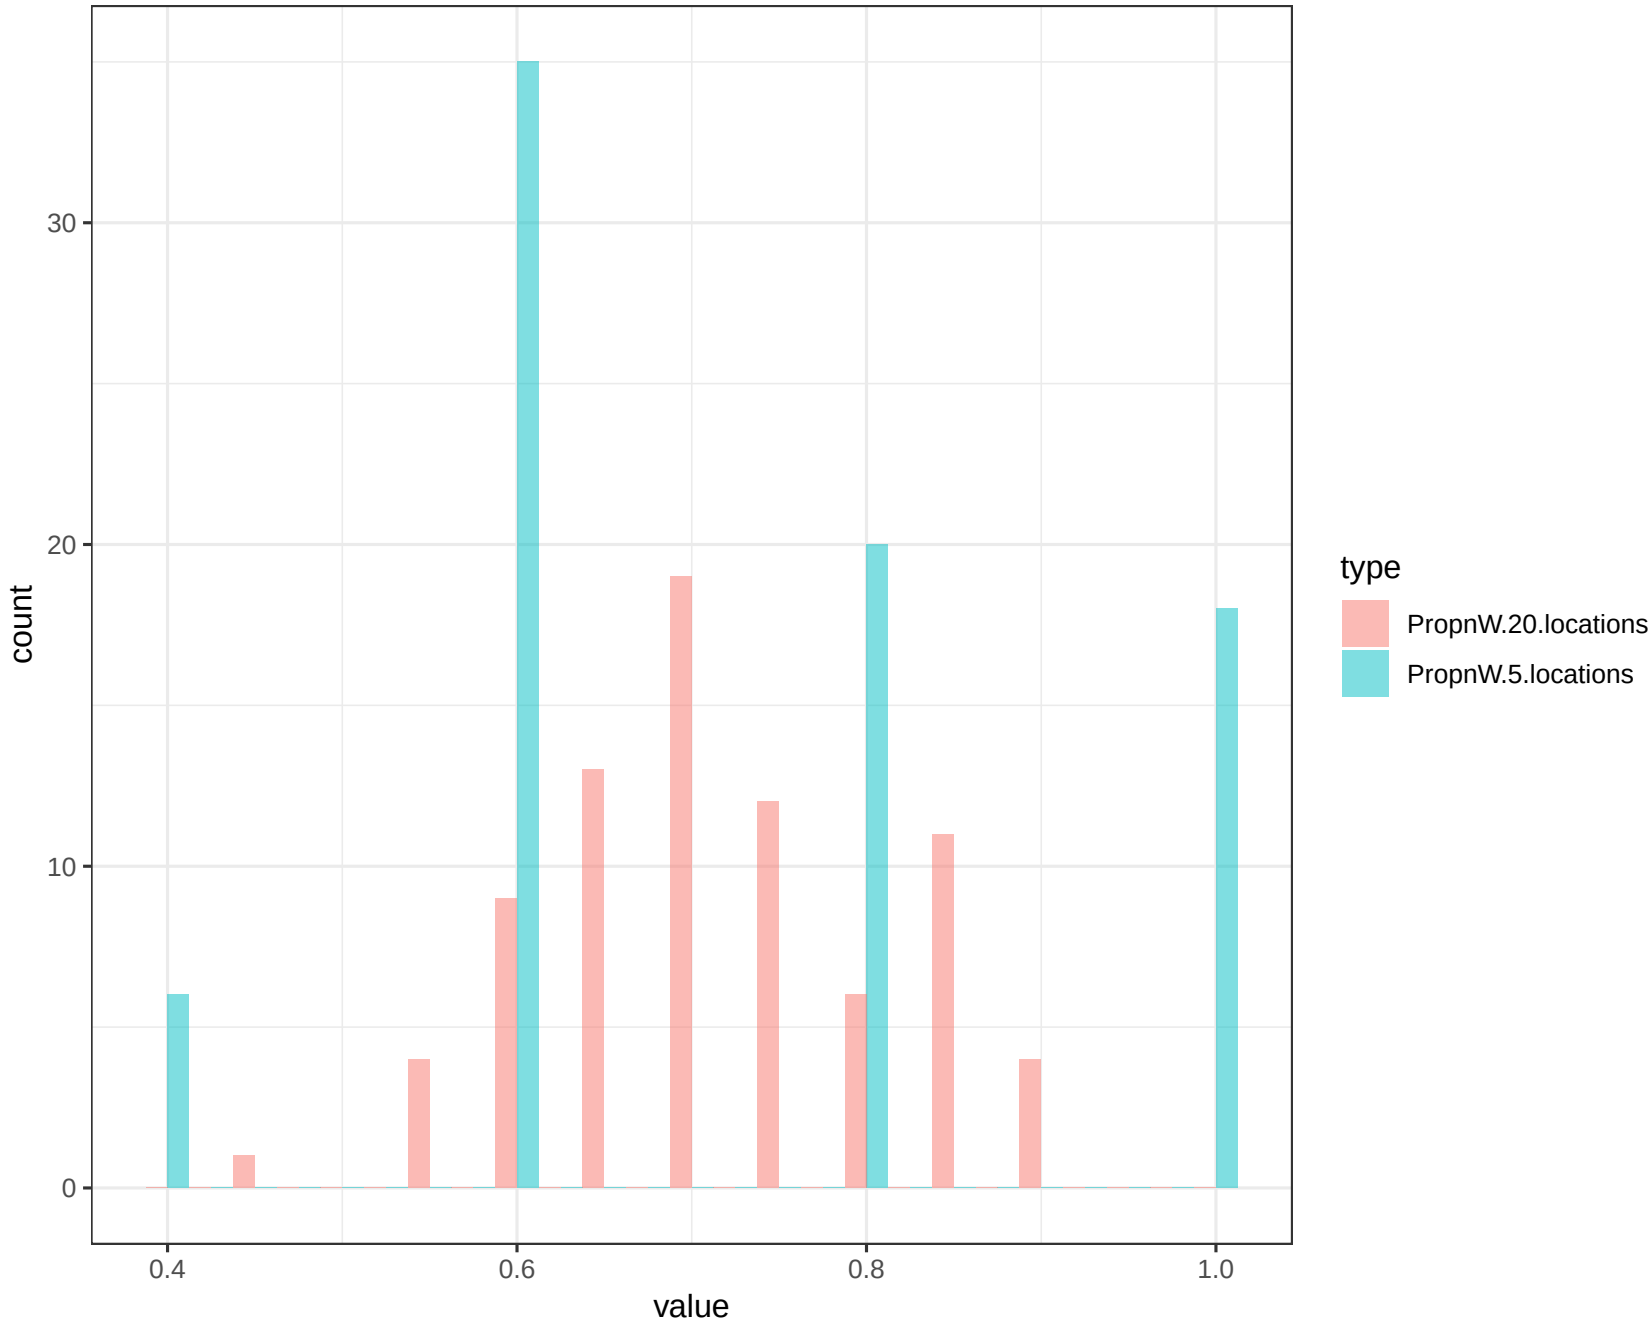

Sampling distribution: mean depth of the ocean

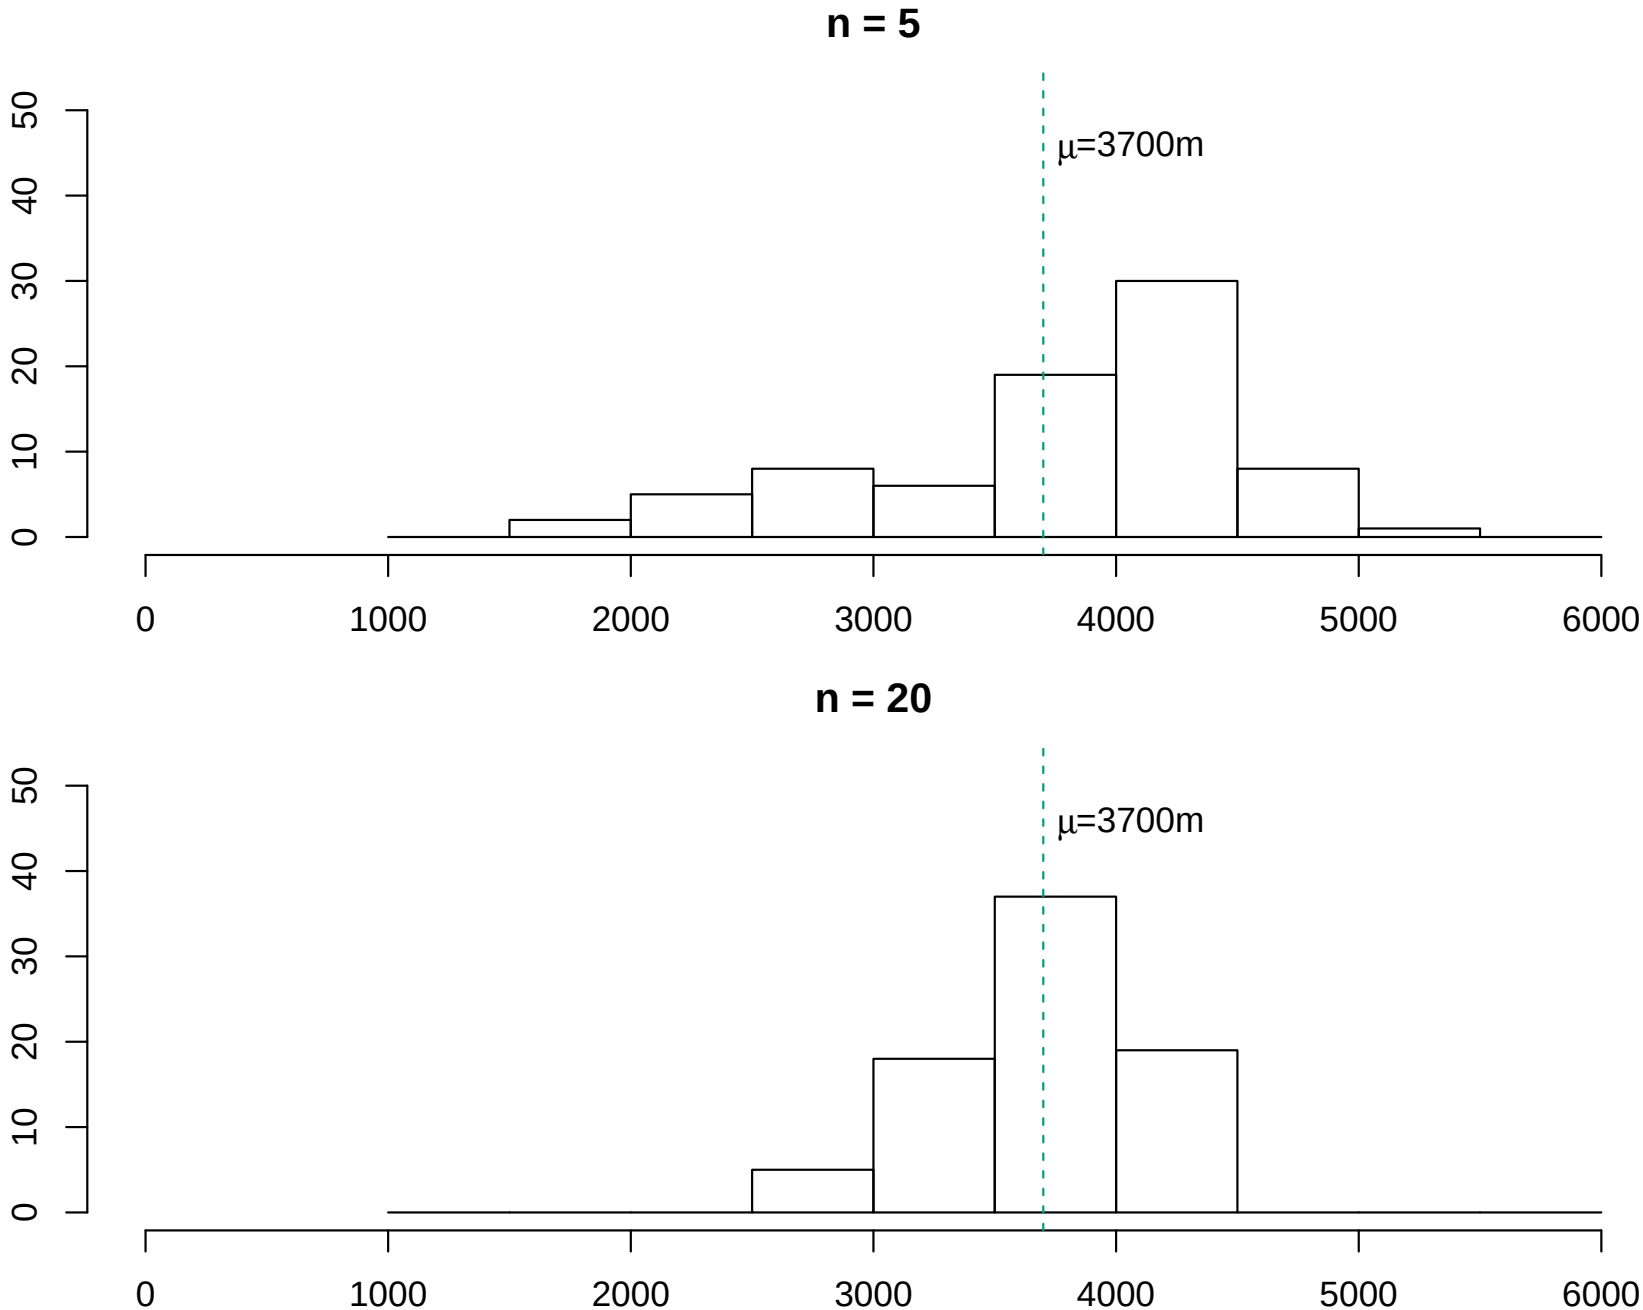

Sampling distribution: mean depth of the ocean

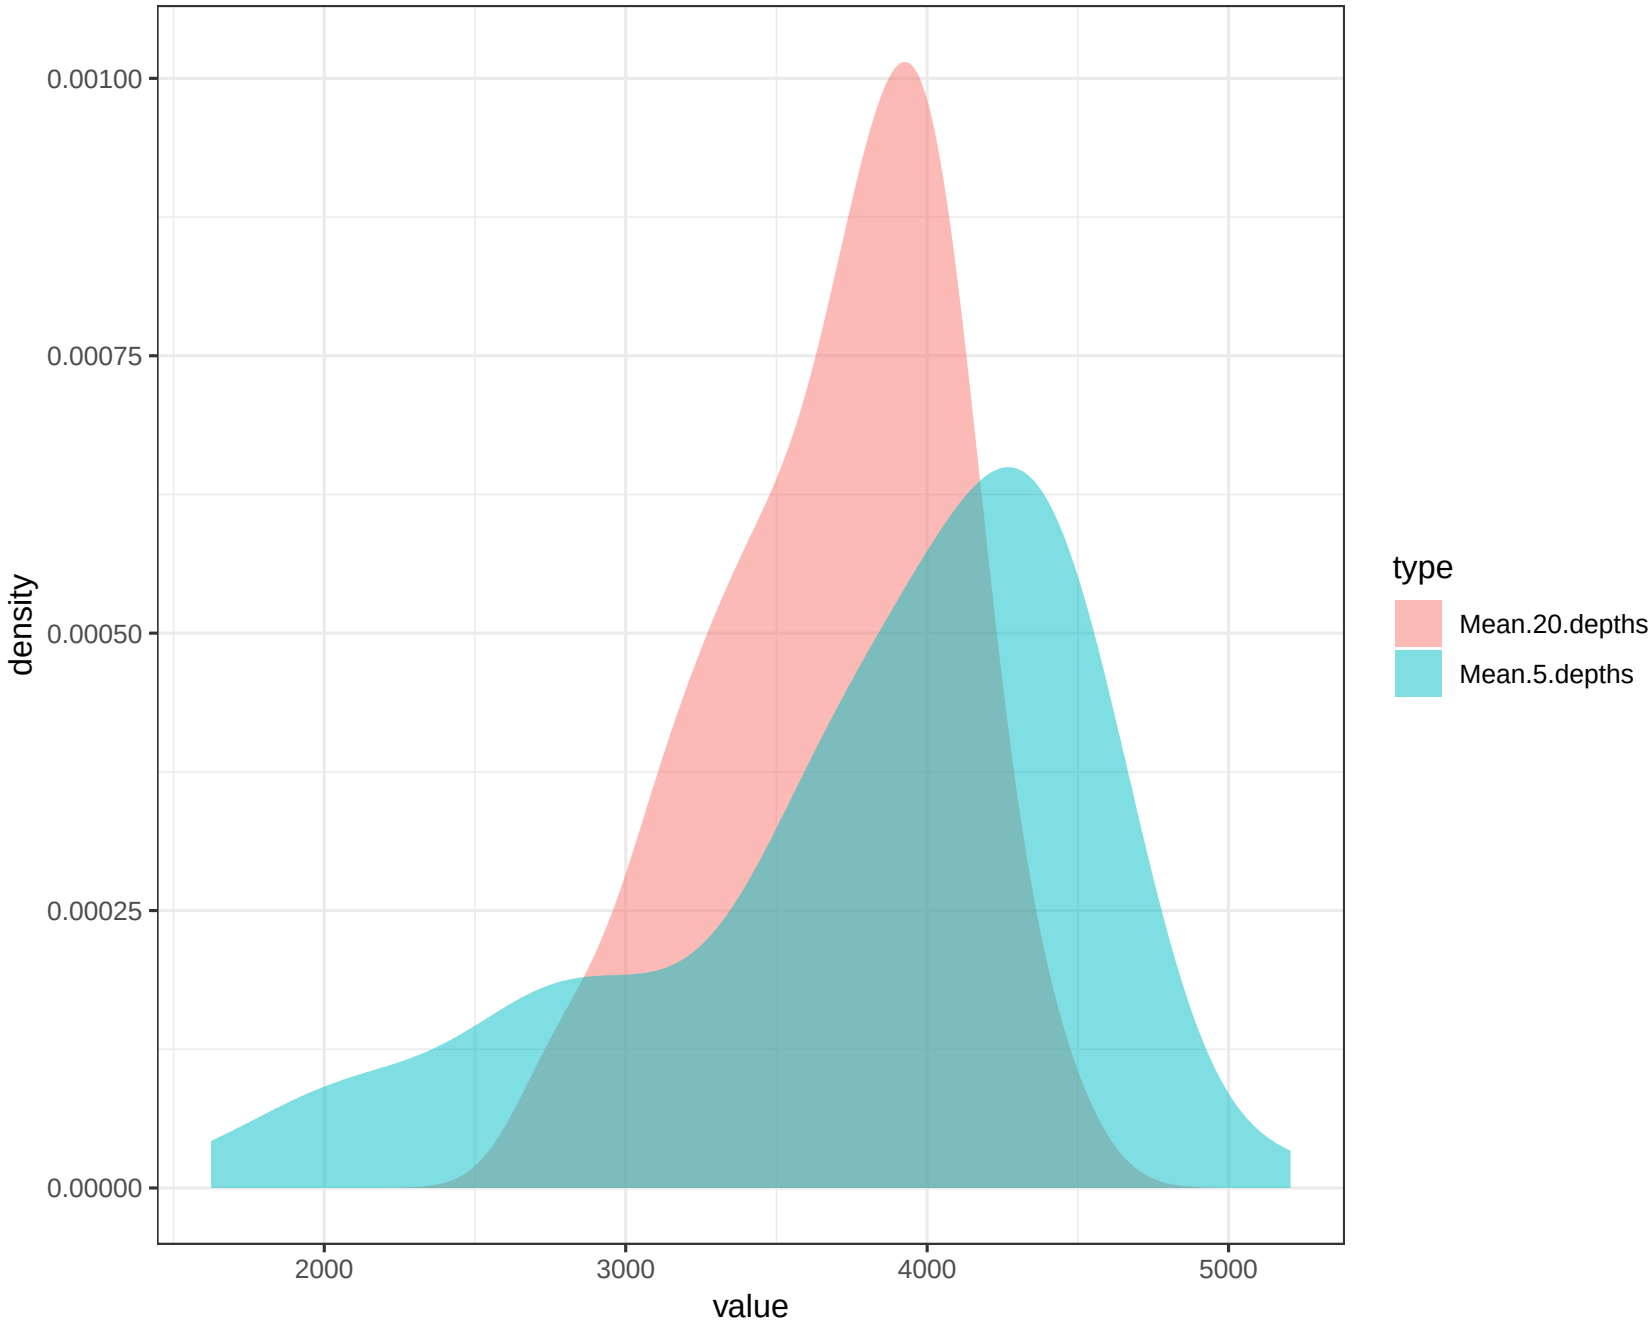

Normal Curves and Calculations

The Normal (Gaussian) distribution

What is it?

- A distribution that describes continuous (numerical) data
- Can also be used to approximate discrete data distributions
- Range is (technically) infinite, though the probability of seeing very large or very small values is extremely tiny
- Fully described by only two parameters, the mean and variance (μ and σ^2)
- NOTE: Baldi & Moore (and R) use the short-hand: $X \sim \mathcal{N}(\mu, \sigma)$, denoting the normal distribution as a function of the mean and *standard deviation*. This is not standard; many texts instead write $X \sim \mathcal{N}(\mu, \sigma^2)$. Be careful of this!

The Normal (Gaussian) distribution

Carl Gauss was a German mathematician who developed a number of important advances in statistics such as the method of least squares.

The Normal distribution

Where do Normal data come from?

- Natural processes
  - Blood pressure
  - Height
  - Weight

- "Man-made" (or derived)
  - Binomial (proportion) and Poisson (count) data are approximately Normal under certain conditions
  - Sums and means of random variables (Central Limit Theorem)
  - Data can sometimes be made to look Normal via transformations (squares, logs, etc)

The Normal distribution

For Normal data, we can use ~~the Gaussian tables~~ R to answer the questions:

- What is the probability that a single observation X is
  - greater than X^*?
  - less than X^*?
  - between X_L^* and X_U^*?
- That is, we can find out information about the percent distribution of X as a function of thresholds X^*, or X_L^* and X_U^*.
- We can also use ~~the Normal tables~~ R to find out information about thresholds X^* that will contain particular percentages of the data. I.e., we can find what threshold values will
  - Exclude the lower $\omega^*\%$ of a population
  - Exclude the upper $\omega^*\%$ of a population
  - Contain the middle $\omega^*\%$ of a population

The Normal distribution

We can use ~~the Gaussian tables~~ **R** to answer these questions **no matter what the values of** μ and σ^2.

That is, the % of the Normal distribution falling between $X_L^* = \mu - m_1\sigma$ and $X_U^* = \mu + m_2\sigma$ where m_1, m_2 are any multiples **remains the same** for any μ and σ.

How so??

Because we can **standardize** any $X \sim \mathcal{N}(\mu, \sigma)$ to find $Z \sim \mathcal{N}(0, 1)$

The Normal distribution

An illustration using IQ scores, which we presume have a $\mathcal{N}(100, 13)$ distribution of scores.

Q1: What percentage of scores are **above** 130?
Two steps:

1. Change of location from $\mu_X = 100$ to $\mu_Z = 0$
2. Change of scale from $\sigma_X = 13$ to $\sigma_Z = 1$

Together, this gives us

$$Z = \frac{X - \mu_X}{\sigma_X} = \frac{130 - 100}{13} = 2.31$$

The Normal distribution

The position of X=130 in a $\mathcal{N}(100, 13)$ distribution is the same as the place of $Z = 2.31$ on the $\mathcal{N}(0, 1)$, which we call the **standardized** Normal distribution (or Z-distribution).

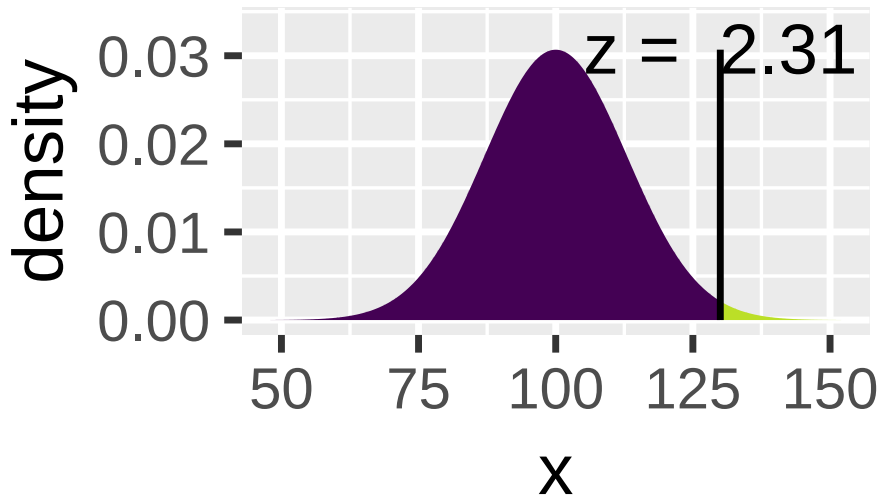

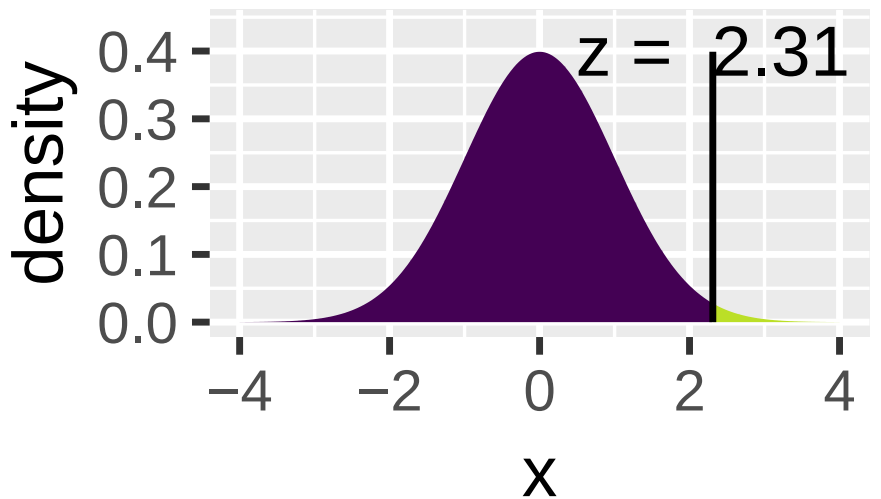

The Normal distribution

How are the values in the Normal tables found?

Normal density:

$$f(x|\mu,\sigma) = \frac{1}{\sqrt{2\pi}\sigma} \exp \frac{-(x-\mu)^2}{2\sigma^2}$$

Probabilities found by integration (area under the Normal curve):

$$P(a \leq x \leq b) = \int_a^b \frac{1}{\sqrt{2\pi}\sigma} \exp \frac{-(x-\mu)^2}{2\sigma^2} dx$$

The Normal distribution

(The percent above $X = 130$) = (% above $Z = 2.31$) =1.04%

How do we know this? We look at the lower tail probability of 2.31 [i.e., the % below 2.31], and then subtract it from 1:

1. $P(X < 130) = P(Z < 2.31) = 0.9896$
2. $P(X > 130) = 1 - P(X < 130) = 0.0104$

So 130 is the 98.96^{th} percentile of a $\mathcal{N}(100,13)$ distribution.

Reminder about percentiles and quantiles

- **Quantile**
  - Any set of data, arranged in ascending or descending order, can be divided into various parts, also known as partitions or subsets, regulated by quantiles.
  - Quantile is a generic term for those values that divide the set into partitions of size n, so that each part represents $1/n$ of the set.
  - Quantiles are not the partition itself. They are the numbers that define the partition.
  - You can think of them as a sort of numeric boundary.

- **Percentile**
  - Percentiles are quite similar to quantiles: they split your set, but only into two partitions.
  - For a generic kth percentile, the lower partition contains k% of the data, and the upper partition contains the rest of the data, which amounts to 100 - k %, because the total amount of data is 100%.
  - Of course k can be any number between 0 and 100.

More about percentiles and quantiles

- In class, we will find ourselves asking for the quantiles of a distribution.
- Percentiles go from 0 to 100
- Quantiles go from any number to any number
- Percentiles are examples of quantiles and you might find some people use them interchangeably (though this may not always be correct since quantiles can take on any value, positive or negative).
- **In particular**, R uses the term quantiles.
- In the previous example, we saw that $P(Z < 2.31) = 0.9896$. In R, 2.31 is called the quantile .

The Normal distribution

(The percent above $X = 130$) = (% above $Z = 2.31$) =1.04%

But wait!! The standard Normal is symmetric about 0, so we can do this another way... The % **above** 2.31 is equal to the % **below** -2.31:

$$P(X > 130) = P(Z > 2.31)$$
$$\Rightarrow P(Z > 2.31) = P(Z < -2.31)$$
$$\Rightarrow P(X > 130) = P(Z < -2.31) = 0.0104$$

So 130 is the 98.96^{th} percentile of a $\mathcal{N}(130, 13)$ distribution. What is the 1.04^{th} percentile?

Transform from $Z = -2.31$ back to X:

$$X = \sigma Z + \mu = 13(-2.31) + 100 = 69.97.$$

For probabilities we use *pnorm*

```
stats::pnorm(q = 130, mean = 100, sd = 13)
```

```
## [1] 0.9894919
```

```
mosaic::xpnorm(q = 130, mean = 100, sd = 13)
```

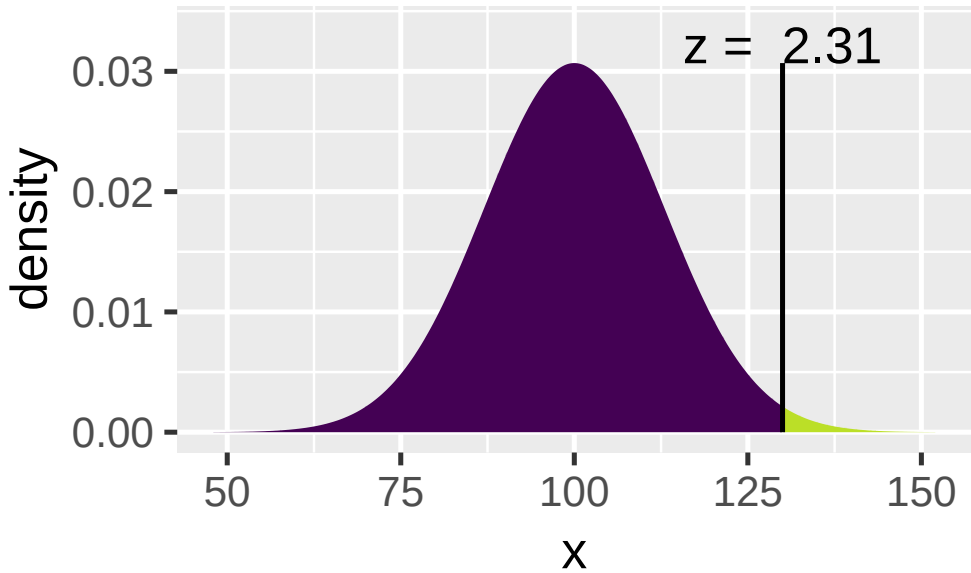


```
## [1] 0.9894919
```

- **pnorm** returns the integral from $-\infty$ to q for a $\mathcal{N}(\mu, \sigma)$
- **pnorm** goes from *quantiles* (think Z scores) to probabilities

For quantiles we use *qnorm*

```
stats::qnorm(p = 0.0104, mean = 100, sd = 13)
```

```
## [1] 69.94926
```

```
mosaic::xqnorm(p = 0.0104, mean = 100, sd = 13)
```

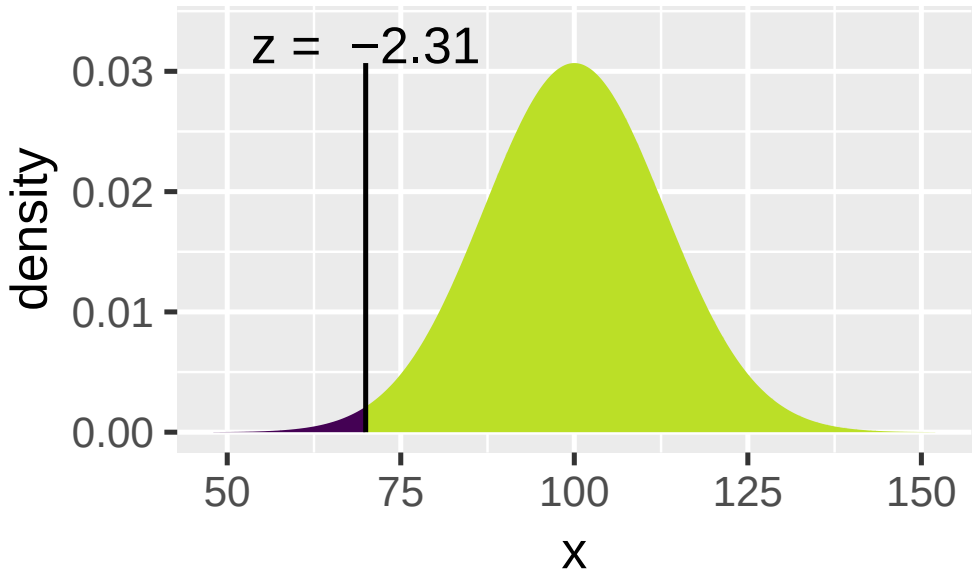


```
## [1] 69.94926
```

- `qnorm` answers the question: What is the Z-score of the *p*th percentile of the normal distribution?
- `qnorm` goes from *probabilities* to quantiles

The Normal distribution

Q2: What is the probability of seeing an IQ score **as extreme as** (think highly unusual) 130?

1. Again, we find that $X = 130$ is the same percentile of the IQ Normal distribution as $Z = 2.31$ is of the standard Normal.
2. To see what scores are as extreme, we want to know the probability that $Z >$2.31 or that $Z <$-2.31.
3. As we saw previously, $P(Z > 2.31) = P(Z < -2.31) = 0.0104$, so the probability of seeing an IQ as extreme or more so than 130 is $2 \times 0.0104 = 0.0208$.

Finding tail probabilities

```
# lower.tail = TRUE is the default
stats::pnorm(q = -2.31, mean = 0, sd = 1, lower.tail = TRUE) +
stats::pnorm(q = 2.31, mean = 0, sd = 1, lower.tail = FALSE)
```

```
## [1] 0.02088815
```

```
mosaic::xpnorm(q = c(-2.31,2.31), mean = 0, sd = 1)
```

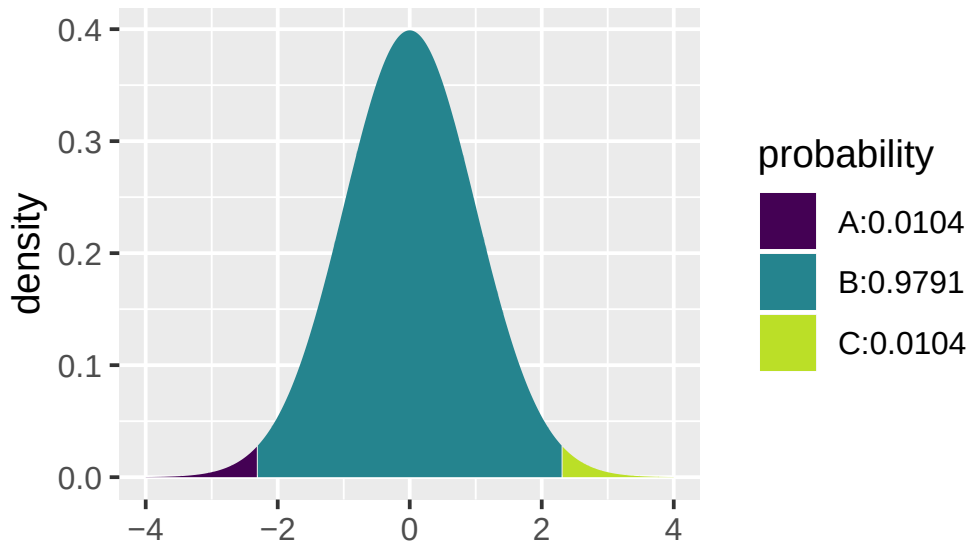


```
## [1] 0.01044408 0.98955592
```

The Normal distribution

Q3: What is the 75^{th} percentile of the IQ scores distribution?
We now have to reverse the sequence of steps:

- Ask yourself: What Z value corresponds to a probability of 0.75? Should you use `pnorm` or `qnorm`?

```
mosaic::xqnorm(p = 0.75, mean = 100, sd = 13)
```

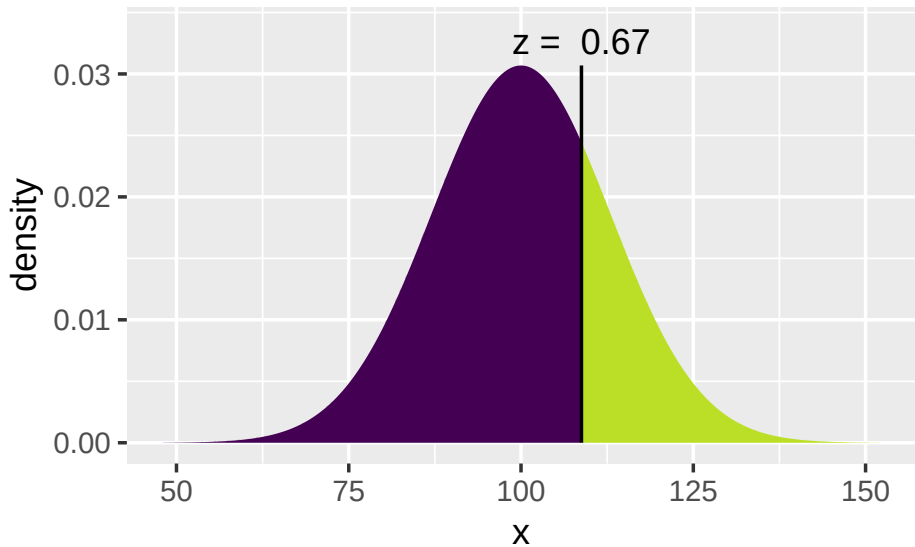


```
## [1] 108.7684
```

This tells us that 75% of the IQ scores fall below 108.8.

Empirical Rule or 68-95-99.7% Rule

In any normal distribution with mean μ and standard deviation σ^2:

- Approximately 68% of the data fall within one standard deviation of the mean.
- Approximately 95% of the data fall within two standard deviations of the mean.
- Approximately 99.7% of the data fall within three standard deviations of the mean.

Demo of Empirical Rule

```
pacman::p_load(mosaic)
pacman::p_load(manipulate)

mNorm <- function(mean = 0, sd = 1) {
  lo <- mean - 5 * sd
  hi <- mean + 5 * sd
  manipulate(
        xpnorm(c(A,B), mean, sd, verbose = FALSE, invisible = TRUE),
        A = slider(lo, hi, initial = mean - sd),
        B = slider(lo, hi, initial = mean + sd)
 )
}
mNorm(mean = 0, sd = 1)
```

Empirical Rule or 68-95-99.7% Rule

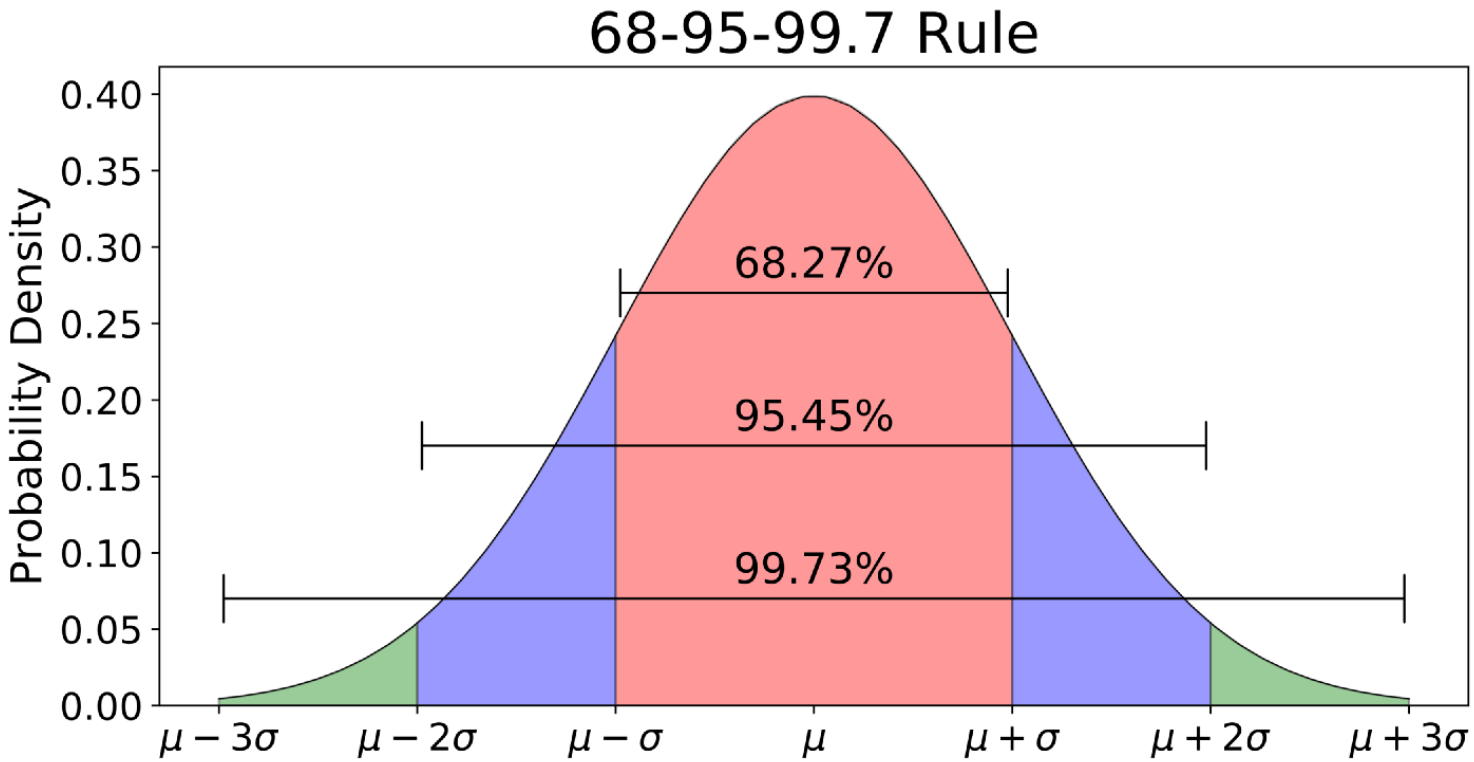

Properties of Normal random variables

Special properties of the Normal distribution:

- In Y is a Normal random variable, then so is $a + bY$.
- If X and Y are two Normal random variables, then $X + Y$ is a Normal random variable. What is the mean and variance of this new random variable?
- If X and Y are two Normal random variables and $\rho_{XY} = 0$ (correlation between X and Y), then X and Y are independent.

Properties of Normal random variables

Let $Y_1, ..., Y_n \sim \mathcal{N}(\mu, \sigma)$, and let each Y_i be independent of the others. (think simple random sample)

Then $\overline{Y} = \frac{1}{n}\sum_{i=1}^{n} Y_i$ has what distribution?

- The sum of Normal random variables is Normal, so $\overline{Y}$ is a Normal random variable.
- $E(\overline{Y}) = \frac{1}{n}\sum_{i=1}^{n} E(Y_i) = \frac{1}{n}\sum_{i=1}^{n} \mu = \mu.$
- $Var(\overline{Y}) = Var(\frac{1}{n}\sum_{i=1}^{n} Y_i) = \frac{1}{n^2}\sum_{i=1}^{n} Var(Y_i) = \sigma^2/n.$
- Standard Error of $\overline{Y}$ = $\sqrt{Var(\overline{Y})} = \sigma/\sqrt{n}$

Central Limit Theorem

Properties of the sample mean: The Central Limit Theorem (CLT)

The sampling distribution of $\overline{Y}$ is Normal if Y is Normal. What probability distribution does the sample mean follow if Y is not Normal?

As sample size increases, the distribution of $\overline{Y}$ becomes closer to a Normal distribution, no matter what the distribution of sampled variable Y!

(This is true as long as the distribution has a finite variance.)

The Central Limit Theorem (CLT)

- The sampling distribution of $\bar{y}$ is, for a large enough n, close to Gaussian in shape no matter what the shape of the distribution of individual Y values.
- This phenomenon is referred to as the CENTRAL LIMIT THEOREM
- The CLT applied also to a sample proportion, slope, correlation, or any other statistic created by aggregation of individual observations

Theorem 1 (Central Limit Theorem)

$$\text{if } Y \sim ???(\mu_Y, \sigma_Y), \text{ then}$$

$$\bar{y} \sim \mathcal{N}(\mu_Y, \sigma_Y/\sqrt{n})$$

Standard error (SE) of a sample statistic

- Recall: When we are talking about the variability of a **statistic**, we use the term **standard error** (not standard deviation). The standard error of the sample mean is $\sigma/\sqrt{n}$.

Remark 1 (SE vs. SD)

In quantifying the instability of the sample mean ($\bar{y}$) statistic, we talk of SE of the mean (SEM)

SE($\bar{y}$) describes how far $\bar{y}$ could (typically) deviate from μ;

SD(y) describes how far an individual y (typically) deviates from μ (or from $\bar{y}$).

CLT in action: Binomial(n = 5,p = 0.8) distribution

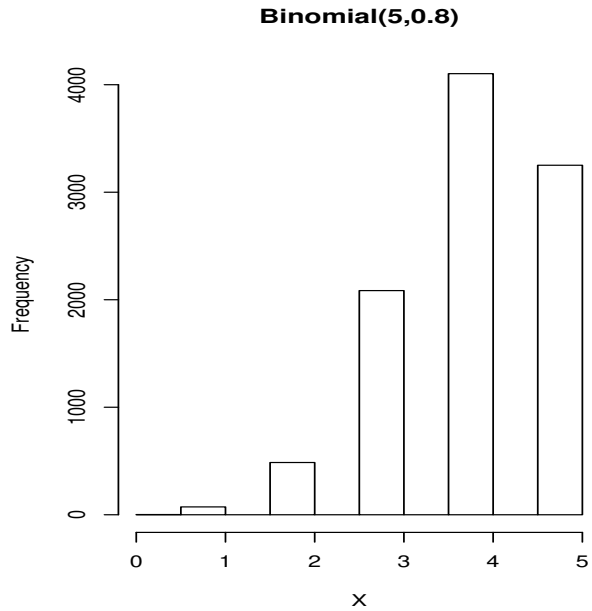

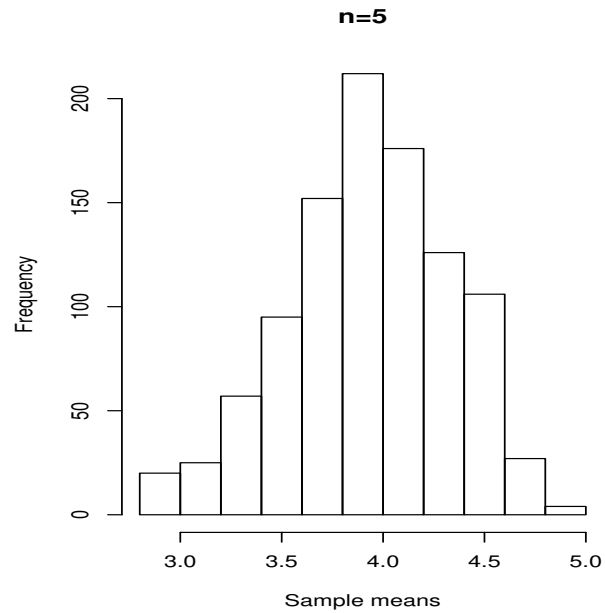

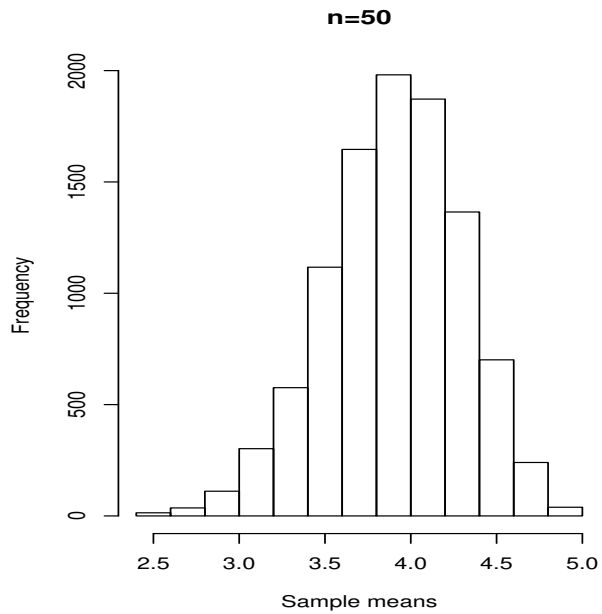

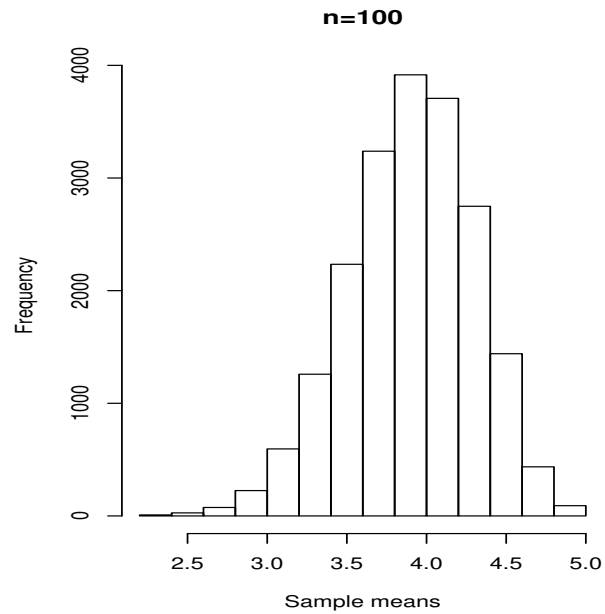

CLT in action: Uniform(a = 0, b = 1) distribution

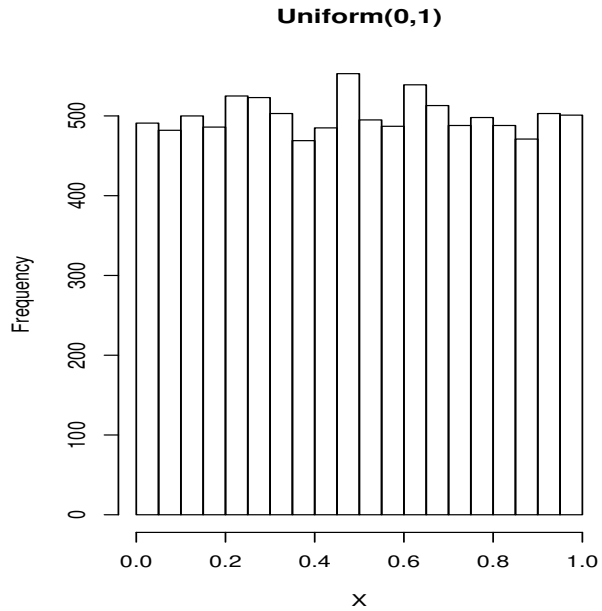

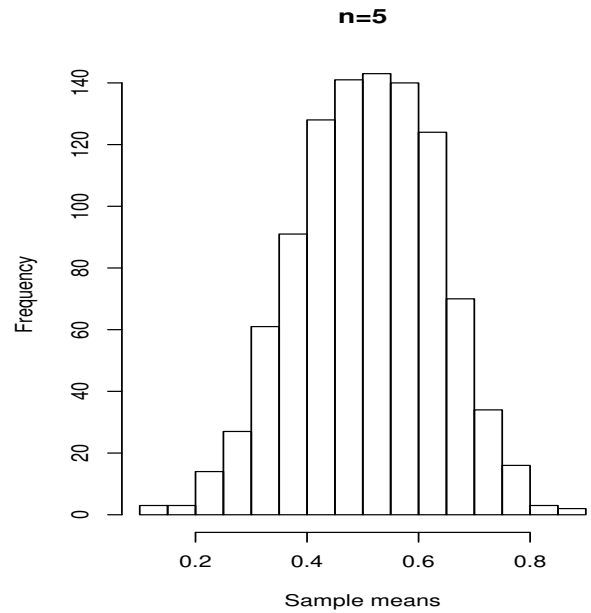

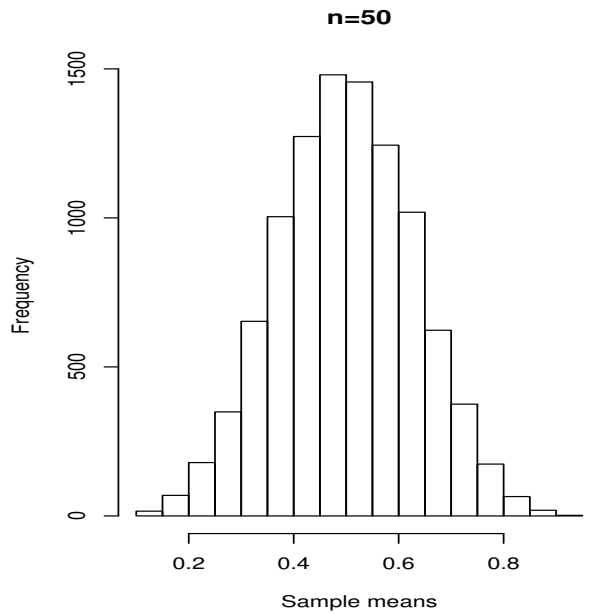

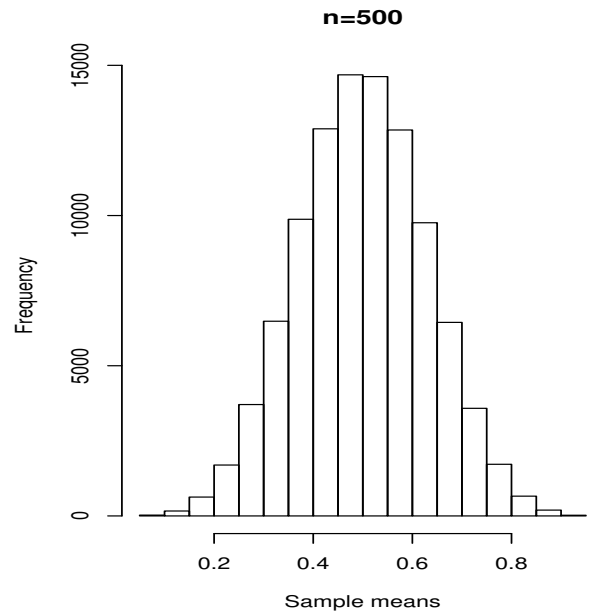

CLT in action: Exponential($\lambda = 1$) distribution

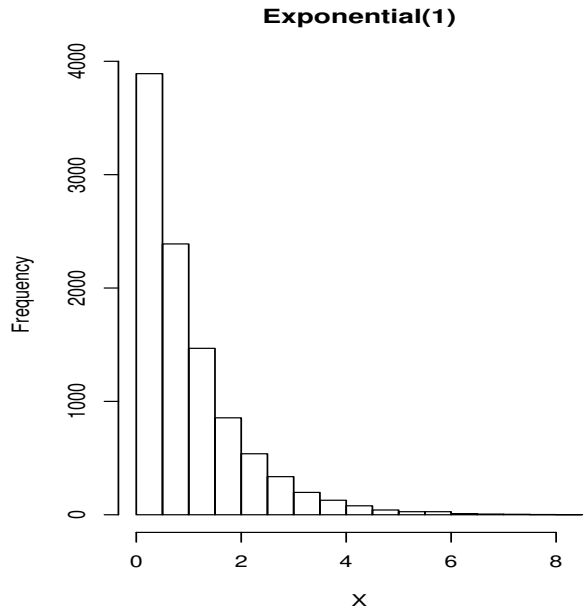

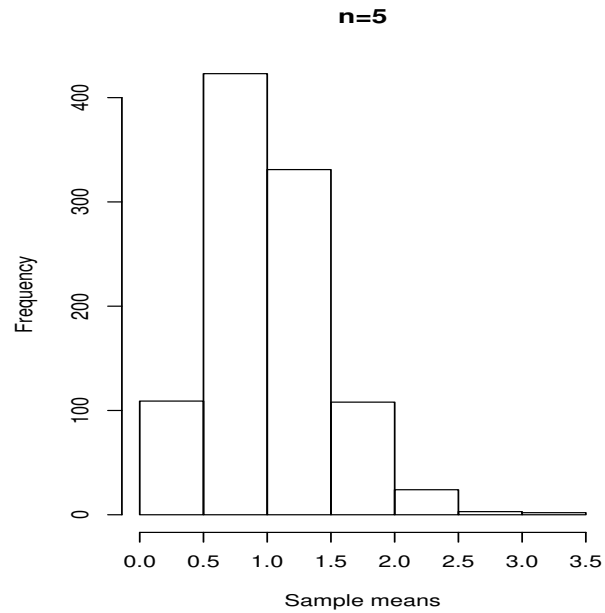

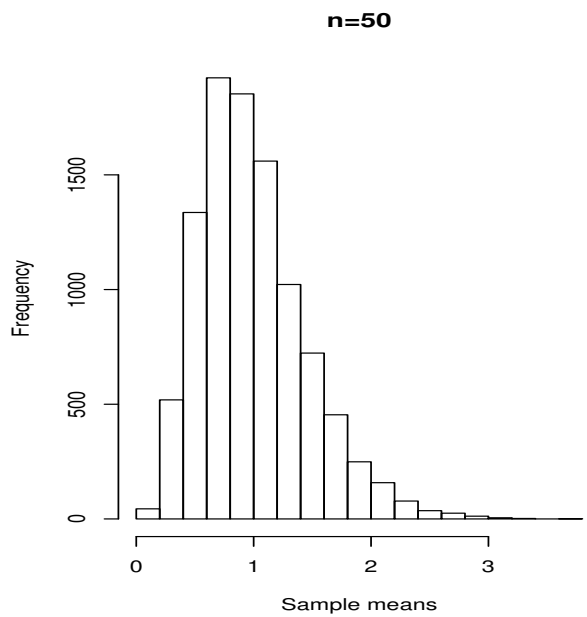

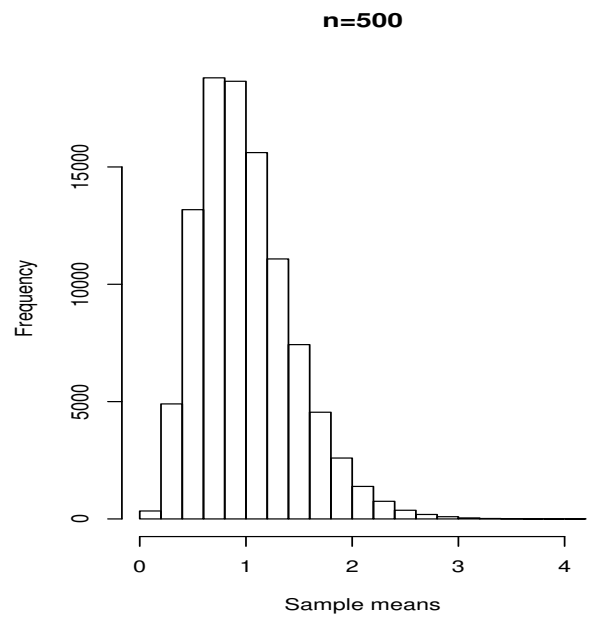

CLT in action: Depths of the ocean

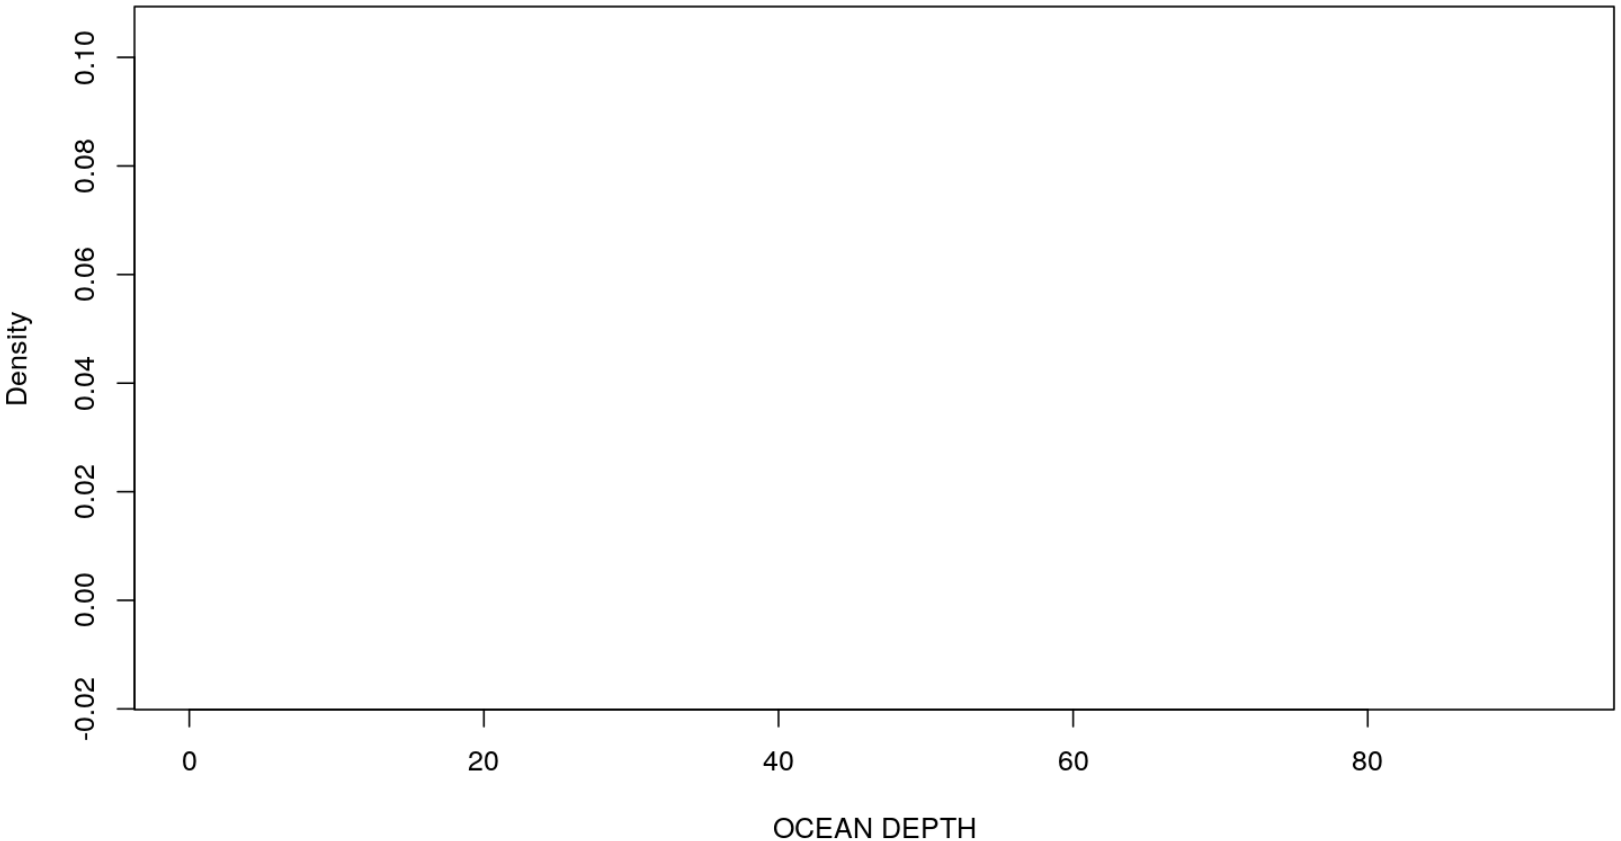

How long does it take for the CLT to 'kick in'?

- How *fast* or slowly the CLT will kick in is a function of how symmetric, or how asymmetric and CLT-unfriendly, the distribution of Y (the depths of the ocean) is

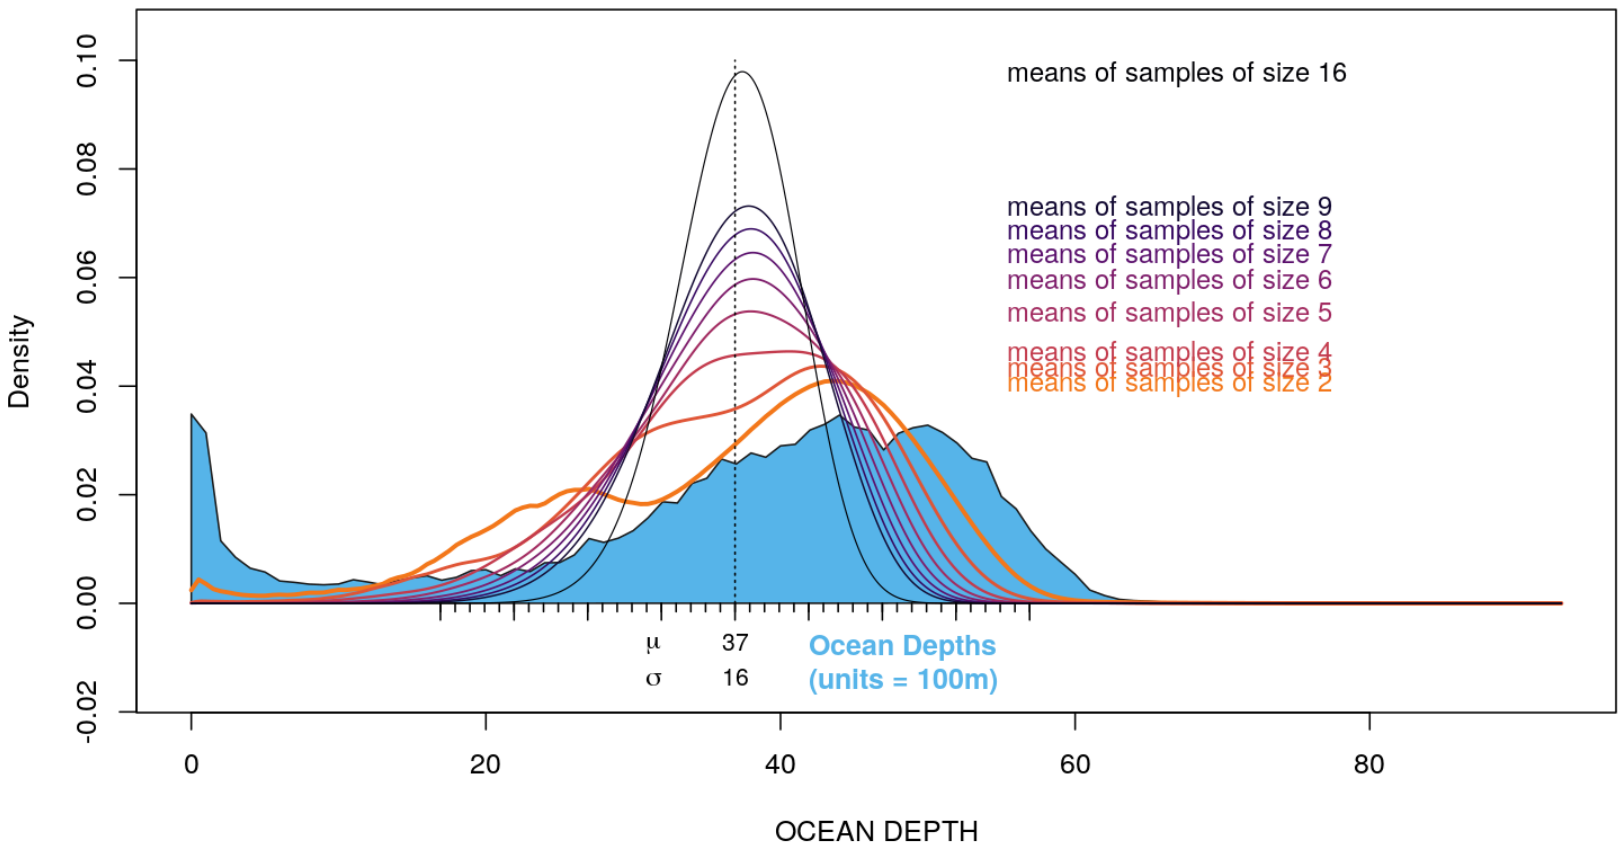

Quadruple the work, half the benefit

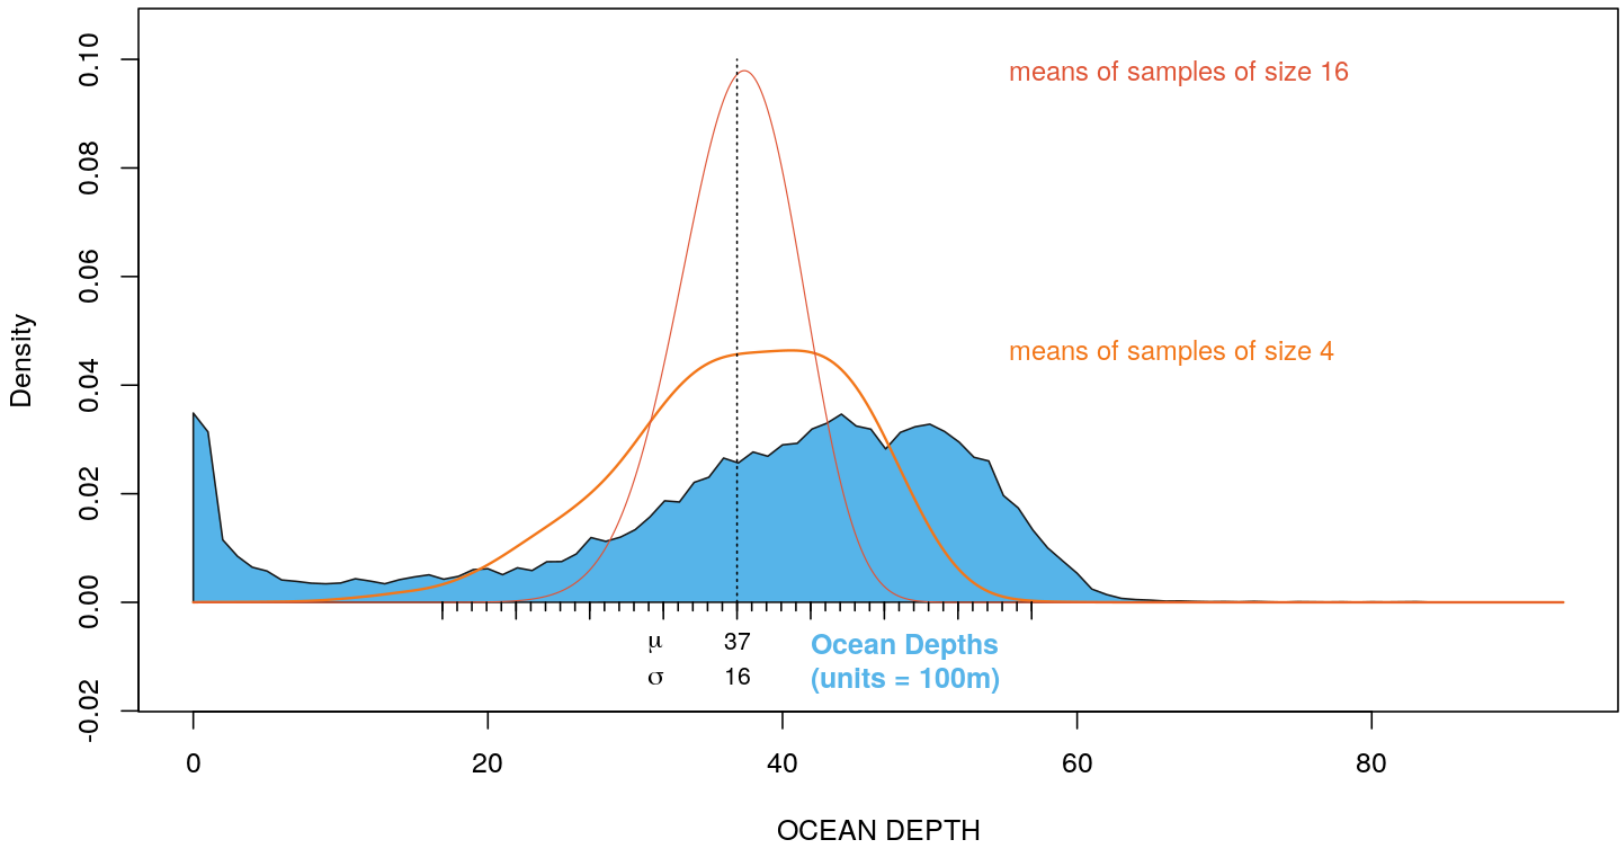

Fig.: When the sample size increases from 4 to 16, the spread of the sampling distribution for the mean is reduced by a half, i.e., the range is cut in half. This is known as the curse of the $\sqrt{n}$

Confidence Intervals

Key takeaways and next steps

1. We've been exclusively talking about point estimates
2. How confident are we about these point estimates?
3. Thought experiment: Estimate the average temperature in Montreal in August over the past 100 years.
4. We're going into stat territory now.

Confidence Interval

Definition 1 (Confidence Interval)

A level C confidence interval for a parameter has two parts:

1. *An interval calculated from the data, usually of the form*

$$estimate \pm margin\ of\ error$$

   where the estimate is a sample statistic and the margin of error represents the accuracy of our guess for the parameter.
2. *A confidence level C, which gives the probability that the interval will capture the true parameter value in different possible samples. That is, the confidence level is the success rate for the method*

Confidence Interval: A simulation study

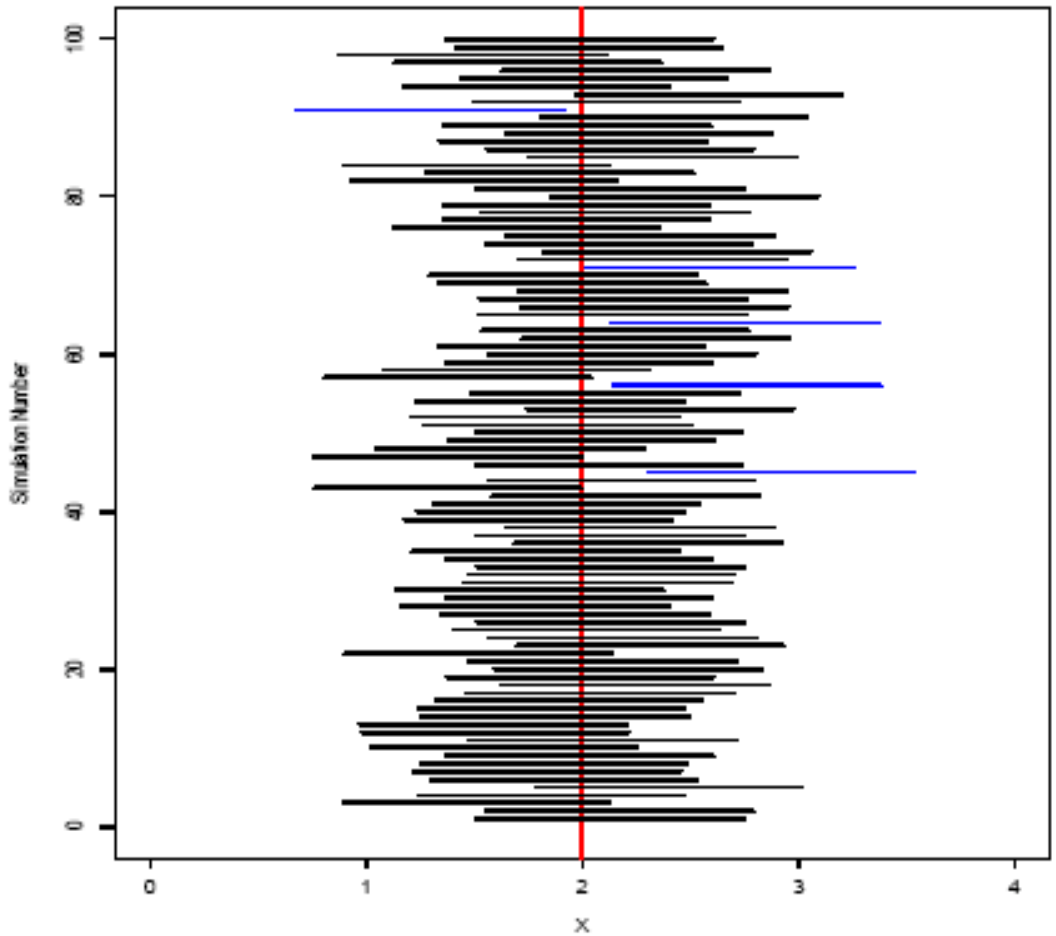

Fig.: True parameter value is 2 (red line). Each horizontal black line represents a 95% CI from a sample and contains the true parameter value. The blue CIs do not contain the true parameter value. 95% of all samples give an interval that contains the population parameter.

Confidence Intervals: we only get one shot

- In practice, we don't take many simple random samples ("repeated" samples) to estimate the population parameter θ.

- Because the method has a 95% success rate, all we need is one simple random sample to compute one CI.

Interpreting a frequentist confidence interval

- The confidence level is the success rate of the method that produces the interval.
- We don't know whether the 95% confidence interval from a particular sample is one of the 95% that capture θ (the unknown population parameter), or one of the unlucky 5% that miss.
- To say that we are 95% confident that the unknown value of θ lies between U and L is shorthand for "We got these numbers using a method that gives correct results 95% of the time."

More about a frequentist confidence interval

- The confidence level of 95% has to say something about the sampling procedure:
  - The confidence interval depends on the sample. If the sample had come out differently, the confidence interval would have been different.
  - With some samples, the interval 'estimate $\pm$ margin of error' does trap the population parameter (the word statisticians use is cover). But with other samples, the interval fails to cover.
- It's like buying a used car. Sometimes you get a lemon – a confidence interval which doesn't cover the parameter.

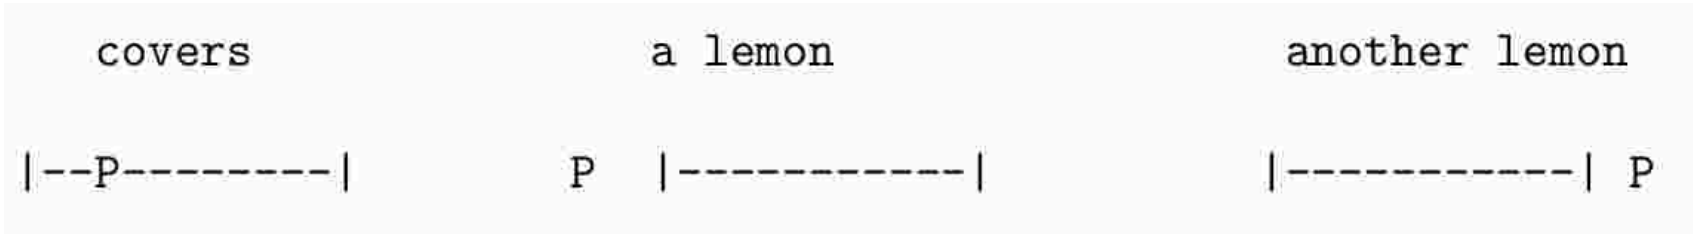

Fig.: 3 confidence intervals 'chasing' (taking a shot at) the population parameter P

More about a frequentist confidence interval

- In the frequentist approach, θ is regarded as a fixed (but unknowable) constant, such as the exact speed of light to an infinite number of digits, or the exact mean depth of the ocean at a given point in time.

- It doesn't "fall" or "vary around" any particular values; in contrast you can think of the statistic $\hat{\theta}$ "falling" or "varying around" the fixed (but unknowable) value of θ

Polling companies

- Polling companies who say “polls of this size are accurate to within so many percentage points 19 times out of 20” are being statistically correct $\rightarrow$ they emphasize the **procedure** rather than what has happened in this specific instance.

- Polling companies (or reporters) who say “this poll is accurate .. 19 times out of 20” are talking statistical nonsense – this specific poll is either right or wrong. On average 19 polls out of 20 are “correct”. But this poll cannot be right on average 19 times out of 20.

Example: Inference for a single population mean

We begin with the (unrealistic) assumption that the population variance is known.

- Then the true variance of the sample mean is known!
- We can use `mosaic::xpnorm(q = c(-1.96, 1.96))` to find that there is a 95% chance that a $\mathcal{N}(0,1)$ random variable lies within 1.96 standard errors of the population mean of the distribution. So then:

$$P\left(-1.96 \leq \frac{\bar{y} - \mu}{\sigma/\sqrt{n}} \leq 1.96\right) = 0.95$$

What does allow us to learn about μ?

Example: Inference for a single population mean

We can use this probability statement about the standardized version of $\bar{y}$ to place bounds on where we think the true mean lies by examining the probability that $\bar{y}$ is within $1.96\frac{\sigma}{\sqrt{n}}$ of μ.

$$
\begin{aligned}
& P\left(-1.96 \leq \tfrac{\bar{y}-\mu}{\sigma/\sqrt{n}} \leq 1.96\right) \\
&= P\left(-1.96\tfrac{\sigma}{\sqrt{n}} \leq \bar{y}-\mu \leq +1.96\tfrac{\sigma}{\sqrt{n}}\right) \\
&= P\left(-\bar{y}-1.96\tfrac{\sigma}{\sqrt{n}} \leq -\mu \leq -\bar{y}+1.96\tfrac{\sigma}{\sqrt{n}}\right) \\
&= P\left(\bar{y}+1.96\tfrac{\sigma}{\sqrt{n}} \geq \mu \geq \bar{y}-1.96\tfrac{\sigma}{\sqrt{n}}\right) \\
&= P\left(\bar{y}-1.96\tfrac{\sigma}{\sqrt{n}} \leq \mu \leq \bar{y}+1.96\tfrac{\sigma}{\sqrt{n}}\right) \\
&= 0.95
\end{aligned}
$$

We call the interval $\left(\bar{y}-1.96\frac{\sigma}{\sqrt{n}}, \bar{y}+1.96\frac{\sigma}{\sqrt{n}}\right)$ a **95% confidence interval** for μ.

Example: Inference for a single population mean

So what does the CI allow us to learn about μ??

- In classical (frequentist) statistics, we assume that the population mean, μ is a **fixed** but unknown value.
- With this view, it doesn't make sense to think of μ as having a distribution. Therefore we can't make probability statements about μ.
- What about the CI? It is made up of the sample mean and other fixed numbers (1.96, the square root of the known sample size n, and the known standard deviation, σ).
- The CI is a random quantity.
- Remember: a random quantity is one in which the outcome is not known ahead of time. We don't know the lower and upper limits of the CI before the sample has been collected since we don't yet know the value of the random quantity $\bar{x}$.

Example: Inference for a single population mean

So what does the CI allow us to learn about μ??

- It tells us that if we repeated this procedure again and again (collecting a sample mean, and constructing a 95% CI), 95% of the time, the CI would *cover* μ.

- That is, with 95% probability, the *procedure* will include the true value of μ. Note that we are making a probability statement about the CI, not about the parameter.

- Unfortunately, we do not know whether the true value of μ is contained in the CI in the particular experiment that we have performed.

Interactive visualization of CIs

http://rpsychologist.com/d3/CI/

Exercise: How deep is the ocean?

1. For your samples of $n = 5$ and $n = 20$ of depths of the ocean, calculate the
   1.1 sample mean ($\bar{y}$)
   1.2 standard error of the sample mean ($SE_{\bar{y}}$)
2. Calculate the 68%, 95% and 99% confidence intervals (CI) for both samples of $n = 5$ and $n = 20$.
3. Enter your results in the Google sheet
4. Plot the CIs for each student using the following code:

```
plot(dt$Mean.5.depths, 1:nrow(dt), pch=20,
xlim=range(pretty(c(dt$lower.mean.5.66, dt$upper.mean.5.66))),
xlab='Depth of ocean (m)', ylab='Student (sample)',
las=1, cex.axis=0.8, cex=1.5)
abline(v = 3700, lty = 2, col = "red", lwd = 2)
segments(x0 = dt$lower.mean.5.66, x1=dt$upper.mean.5.66,
y0 = 1:nrow(dt), lend=1)
```